

# Deep Learning For Financial Time Series

---

Xiaodong YANG  
[yangxiaodong993993@gmail.com](mailto:yangxiaodong993993@gmail.com)

Final Report  
Certificate in Quantitative Finance  
June 2022 Program



CQF | INSTITUTE

January 2023

# Abstract

The direction of the financial market is always stochastic and volatile but is important for financial market anticipates. Analysts now are trying to apply the modelling techniques from Natural Language Processing in the field of Finance to the similarity of having the sequential property in the data. In this research, I have constructed and applied the state-of-art deep learning sequential model, namely Long Short Term Memory Model (LSTM), Stacked-LSTM and LSTM with optimised hyperparameter into the prediction of stock prices on the next day. Moreover, using the prediction, I built up two trading strategies and compared them with the benchmark. My input data contains all technical indicators, which are carefully selected by feature engineering and applied to LSTM models. The result has shown that LSTM with an optimised hyperparameter beats all other models in terms of prediction error and shows much higher returns in my trading strategy over other models. Furthermore, I discovered that the stacked-LSTM model can improve the predictive power over the single LSTM since it has more complex model structures.

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Data Processing</b>	<b>2</b>
2.1	Data Collection . . . . .	2
2.2	Technical Indicators . . . . .	2
2.3	Labels . . . . .	2
2.4	Imbalanced Class . . . . .	3
<b>3</b>	<b>Feature Engineering</b>	<b>4</b>
3.1	Boruta . . . . .	4
3.2	K-Means Clustering . . . . .	4
3.3	Self Organising Maps . . . . .	6
<b>4</b>	<b>Exploratory Data Analysis</b>	<b>7</b>
4.1	Features Describe . . . . .	8
4.2	Correlation Matrix . . . . .	8
4.3	Pairwise Plot . . . . .	9
<b>5</b>	<b>Experiments And Results</b>	<b>10</b>
5.1	Simple LSTM . . . . .	11
5.2	Stacked-LSTM Model . . . . .	12
5.3	Hyperparameter Optimization LSTM . . . . .	14
<b>6</b>	<b>Trading Strategies</b>	<b>16</b>
6.1	Long-Only Strategy . . . . .	16
6.2	Long-Short Strategy . . . . .	17
<b>7</b>	<b>Conclusion</b>	<b>18</b>
<b>Appendix</b>		<b>19</b>
Python Code . . . . .		19
<b>References</b>		<b>41</b>

# List of Figures

1	SPY Adjusted Close Prices . . . . .	2
2	SPY Daily Returns . . . . .	2
3	Class Distribution . . . . .	3
4	Elbow Plot . . . . .	5
5	Silhouette Score . . . . .	5

6	SPY Features Clustering . . . . .	6
7	Self Organising Maps . . . . .	7
8	Feature Descriptions . . . . .	8
9	Feature Correlation Matrix . . . . .	9
10	Feature Pairwise Plot . . . . .	10
11	Simple LSTM Model . . . . .	11
12	Simple LSTM Confusion Matrix . . . . .	12
13	Simple LSTM ROC . . . . .	12
14	Stacked-LSTM Model . . . . .	13
15	Stacked LSTM Confusion Matrix . . . . .	14
16	Stacked LSTM ROC . . . . .	14
17	Hyperparameter Optimization LSTM . . . . .	15
18	Optimised LSTM Confusion Matrix . . . . .	16
19	Optimised LSTM ROC . . . . .	16
20	Long-Only Strategy Return . . . . .	17
21	Long-Short Strategy Return . . . . .	18

## List of Tables

1	Features List . . . . .	8
2	Long-Only Returns . . . . .	16
3	Long-Short Returns . . . . .	17

## List of Algorithms

1	K-Means Clustering Algorithm . . . . .	5
2	Self Organising Maps Algorithm . . . . .	6

# 1 Introduction

In the field of quantitative trading, predicting the future security returns lies in the center of the industry, as the future trading strategy is always deployed and created based on our view of the financial market in the future. The trading area has two main different methods, namely fundamental analysis and quantitative trading. Fundamental method makes the trading decision based on the subjective view on the industry or company's future direction, it mainly relies on the public information such as market news, corporate statistics as well as a series of financial statement releases. On the other hand, the quantitative trading strategy uses mathematical models to make the decision hence avoids the interruption of human subjectivity and emotion. Traditionally, the methodology of quantitative strategy involves using linear regressions, ARIMA model as well as GARCH model to capture the features of time series and the stochasticity of the volatility. These methods were proved to be effective for a certain period of time in the old regimes. As the regime shift happens in the financial industry, these models became less effective. The quantitative trading industry has therefore shifted into the ‘deep learning era’, which has been more frequently used nowadays. In this research, we predicted stocks return using the deep learning model, more specifically LSTM and Attention-LSTM models. Conventional financial time series prediction only uses price and volume to predict the future direction. In this work, the input includes abundant technical indicators calculated by open, high, low, adjusted close price, volumes, etc. Usually these technical indicators have significant impacts on future price movements. The output of our model is the return directions on the next day. Once the future return direction is predicted, we will build up a quantitative trading strategy based on the prediction. The return of our strategy is compared with the market [Zou & Qu \(2020\)](#).

SPDR S&P 500 ETF Trust is chosen in this project, which is one of the earliest investment vehicles created. The SPY aims to track the S&P500 index as closely as possible, holding stocks in direct proportion to their weights within the index. Although the SPY expense ratio is slightly higher than other ETF’s that also track the S&P500 index, the SPY’s history and popularity make it an interesting and challenging ETF to forecast. Given that the trading volume of the S&P500 is so high amongst both retail and institutional investors, any edge that exists is unlikely to remain for long, being rapidly traded out as soon as it appears. This sides slightly with the efficient market hypothesis to some degree, but there is significantly more achievement and interest embedded in beating one of the most popular and rapidly traded instruments available.

## 2 Data Processing

### 2.1 Data Collection

The data I have for this research was fully available and I obtained it from yahoo finance by finance API. These data were compiled in the period 31/12/2010 to 31/12/2022, which was divided into two subsets: the training set and the test set, respectively with ratios of 70% and 30%. The following will describe the data. Figure 1 and Figure 2 show the SPY performance to date.

Figure 1: SPY Adjusted Close Prices

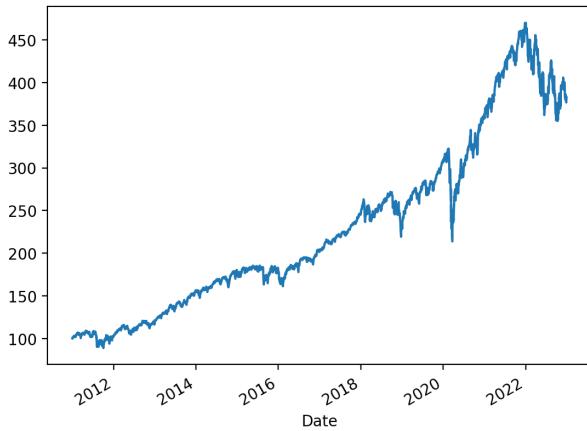
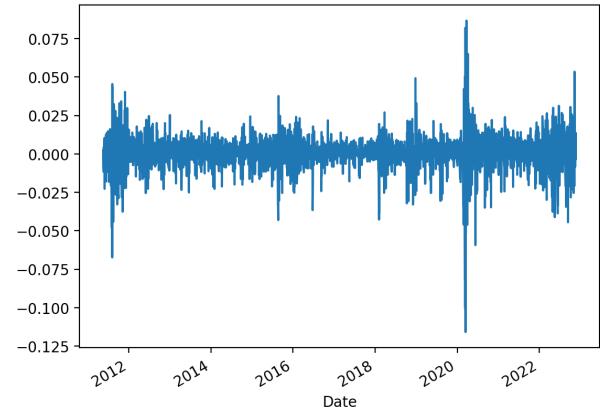


Figure 2: SPY Daily Returns



### 2.2 Technical Indicators

Technical indicators are available from Technical Analysis Library, including Momentum Indicators, Volume Indicators, Volatility Indicators, Trend Indicators, etc. The features are constructed by these technical indicators. Before feature engineering, several irrelevant variables and NA data need to be removed from the features shown by the below code

```
stock_data.drop(['HIL01_13_21', 'HIL0s_13_21', 'PSARl_0.02_0.2',
                 'PSARs_0.02_0.2', 'PSARaf_0.02_0.2', 'QQE1_14_5_4.236',
                 'QQEs_14_5_4.236', 'SUPERTl_7_3.0', 'SUPERTs_7_3.0',
                 'Open', 'High', 'Low', 'Close', 'Adj Close'],
                 axis=1, inplace=True)
stock_data = stock_data.dropna(axis=0)
```

### 2.3 Labels

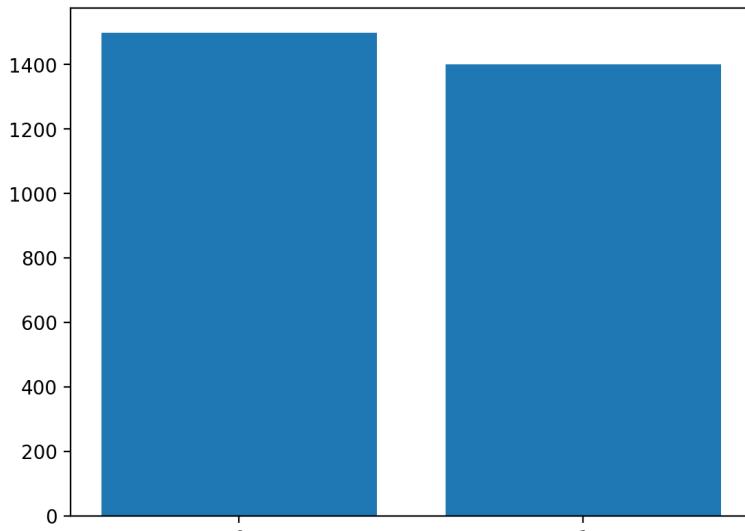
Then "0" and "1" are used to encode the return. To avoid our predictions' accuracy being influenced by the noise that is extremely small near-zero returns, a threshold was set as 0.09%, which is the expense ratio of SPDR S&P 500 ETF Trust referring to [Yahoo Finance](#). In other

words, the returns exceeding its expense ratio were labelled as 1, otherwise, 0, meaning the long indicator will be made only when its return can cover its expense ratio. The labelling criterion is shown by

$$y = \begin{cases} 1, & \text{Returns} > \text{Threshold} \\ 0, & \text{Otherwise} \end{cases}$$

## 2.4 Imbalanced Class

Figure 3: Class Distribution



Source: My Calculations

Another issue impacting the classification negatively is an unbalanced class of the labels, which might lead to situations in which the [learning algorithm] neural network simply predicts the class with the higher frequency since this already can lead to low loss and high accuracy values [Hilpisch \(2020\)](#). Through observing Figure 3, the frequency of '0' is higher than '1', meaning there is an imbalanced class issue. Hence, it can be fixed by applying appropriate weights to make sure that both classes gain equal importance during the DNN training step [Kannan Singaravelu \(2022a\)](#). From the below code, labels '0' and '1' should be applied weights of 0.9670 and 1.0354 respectively.

```
# class frequency
c = stock_data['Returns'].value_counts()

# class weight function
def cwts(stock_data):
    c0, c1 = np.bincount(stock_data['Labels'])
    w0=(1/c0)*(len(stock_data['Returns']))/2
    w1=(1/c1)*(len(stock_data['Returns']))/2
    return {0: w0, 1: w1}
```

```
# check class weights
class_weight = cwts(stock_data)
[Out] {0: 0.9670, 1: 1.0354}
```

## 3 Feature Engineering

The dimension of data should be reduced. Because high-dimension data might cause two issues. It would cause overfitting when all features are used to train the model. Besides, high-dimension might influence computing speed. Therefore, three methods are employed to reduce the dimension of data, including Boruta, K-Means Clustering, and Self-Organising Maps.

### 3.1 Boruta

Boruta algorithm is a wrapper built around the random forest classification algorithm, which is relatively quick, can usually be run without tuning of parameters and gives a numerical estimate of the feature importance. Boruta is based on stochasticity which forms the foundation of the random forest classifier. And it adds randomness to the system and collects results from the ensemble of randomized samples one can reduce the misleading impact of random fluctuations and correlations where this extra randomness shall provide us with a clearer view of which attributes are really important [Kursa & Rudnicki \(2010\)](#).

In my study, the features and target were used to train the classifier and got classification. All features are ranked by specific order based on their relevance. Rank 1 means the most relevance. Hence, all features ranked in '1' were selected for the next step. After that, the feature number is reduced from 272 to 64.

### 3.2 K-Means Clustering

K-Means Clustering is a branch of unsupervised machine learning models that seeks to learn from the properties of the data by identifying groups or clusters in the dataset. The difference between clustering and classification is that the former seeks to identify a set of similar data points and calls the resulting set a cluster [Xu & Tian \(2015\)](#).

The K - Means Clustering algorithm searches for a predetermined number of clusters within an unlabeled dataset. There are two assumptions that the optimal cluster [Kannan Singaravelu \(2022b\)](#)

- cluster center which is the mean of all the points belonging to the cluster
- each point is closer to its own cluster center than to other cluster centers

The K-Means Algorithm seeks to find k clusters within a data set. The clusters are chosen to

reduce the inertia, the objective function

$$\min_{C_k} \sum_k \sum_{X_j \in C_k} \|X_j - \mu_k\|^2$$

and the centroid of a cluster  $\mu_k$  is equal to

$$\mu_k = \frac{1}{C_k} \sum_{X_j \in C_k} X_j$$

Where  $C_k$  is the number of points in cluster  $k$ .

The cluster assignments are done based on squared Euclidean distance and the algorithm tries to minimize the objective function. This process is repeated, iteratively, until the sum of the distance is minimized [Kannan Singaravelu \(2022b\)](#).

---

#### Algorithm 1 K-Means Clustering Algorithm

---

- 1: Randomly selecting  $k$  -centroids or cluster center
  - 2: **while** Not Converge ... **do**
  - 3:     a. E-Step: assign points to the nearest cluster center
  - 4:     b. M-Step: set the cluster center to the mean
  - 5: **end while**
- 

Algorithm 1 was utilised in this part to analyse the clustering. In this algorithm, the number of clusters is a hyperparameter. Hence, it is crucial for the model to select the optimal number of clusters. Too many or few clusters are not appropriate. Instead, The Elbow plot was employed to identify the infection point where increasing the number of clusters no longer results in a significant drop in inertia and select the optimal number of clusters [Kannan Singaravelu \(2022b\)](#).

Figure 4: Elbow Plot

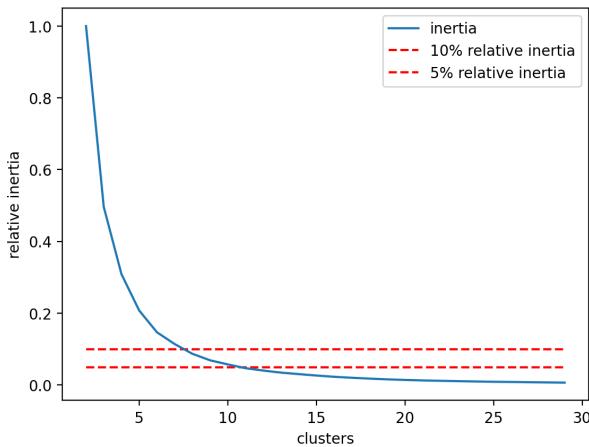
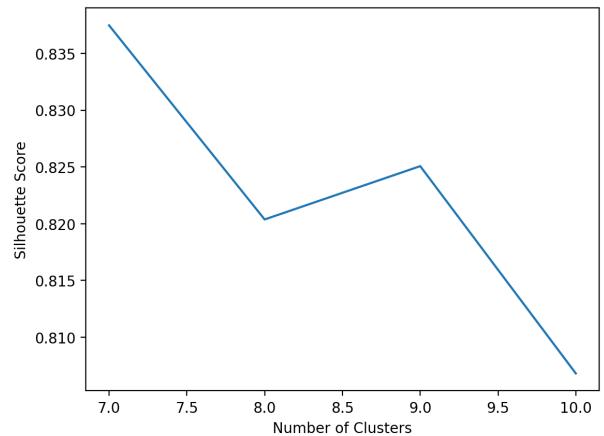


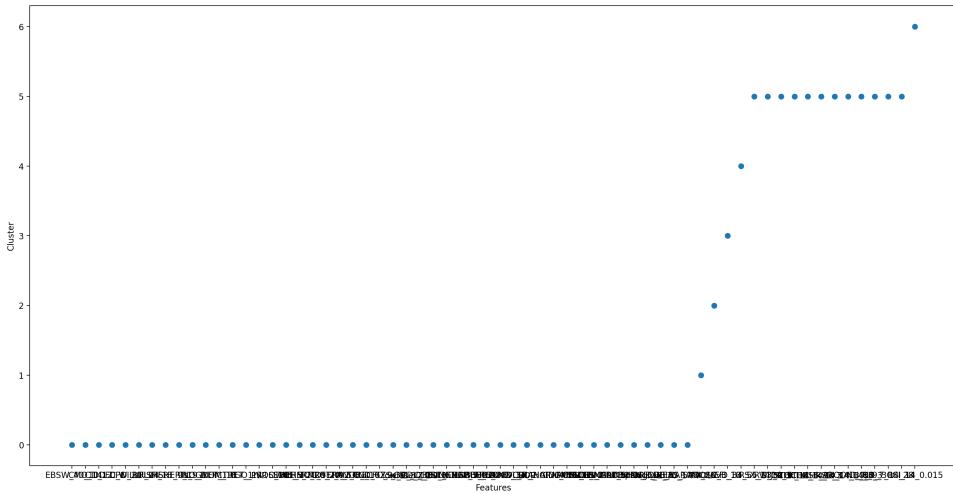
Figure 5: Silhouette Score



The K-Means implementation seeks to find the minimum number of clusters required to achieve low, stable inertia. From Figure 4, it seems the range of the optimal value is from 7 to 10 in

relative inertia from 5% to 10%. Besides, Silhouette Coefficient is another method to measure how dense and separated are the clusters. And its range is from -1 to 1. The samples with larger values are closer to their clusters than others. No doubt that the Silhouette coefficient implementation is more reliable than the Elbow plot since it provides the precise score, which is straightforward Zhou & Gao (2014). Figure 5 shows the Silhouette coefficients with the range of cluster numbers from 7 to 10. It includes that a cluster number equal to 7 has the highest Silhouette Score, meaning 7 is the optimal cluster number in this algorithm. From Figure 6, most of the features are at cluster "5".

Figure 6: SPY Features Clustering



Source: My Calculations

### 3.3 Self Organising Maps

The self-organising map (SOM) is one of the widely implemented neural network models, which is an unsupervised learning method and computationally simple. It produces a topology-preserving mapping between high-dimensional input space and low-dimensional map space. Hence, it is usually used for feature detection [Lau et al. \(2006\)](#). SOM applies competitive learning when compared to error-correlated learning, which involves backpropagation and gradient descent. In SOM, the training data has no labels and the map learns to differentiate and distinguish features based on similarities [Kannan Singaravelu \(2022d\)](#). The SOM algorithm is shown by Algorithm 2.

**Algorithm 2** Self Organising Maps Algorithm

- 1: Initialise weights  $w_{ij}$
  - 2: Calculate distance  $d_{ij} = \sum(w_{ij} - x_i)^2$
  - 3: Calculate  $w_{ij}(t + 1) = w_{ij}(t) + \theta(t) \cdot \alpha(t) \cdot (x_i(t) - w_{ij}(t))$

Where

$x_i(t)$  is the characteristics of the samples

$\alpha(t)$  is a learning rate

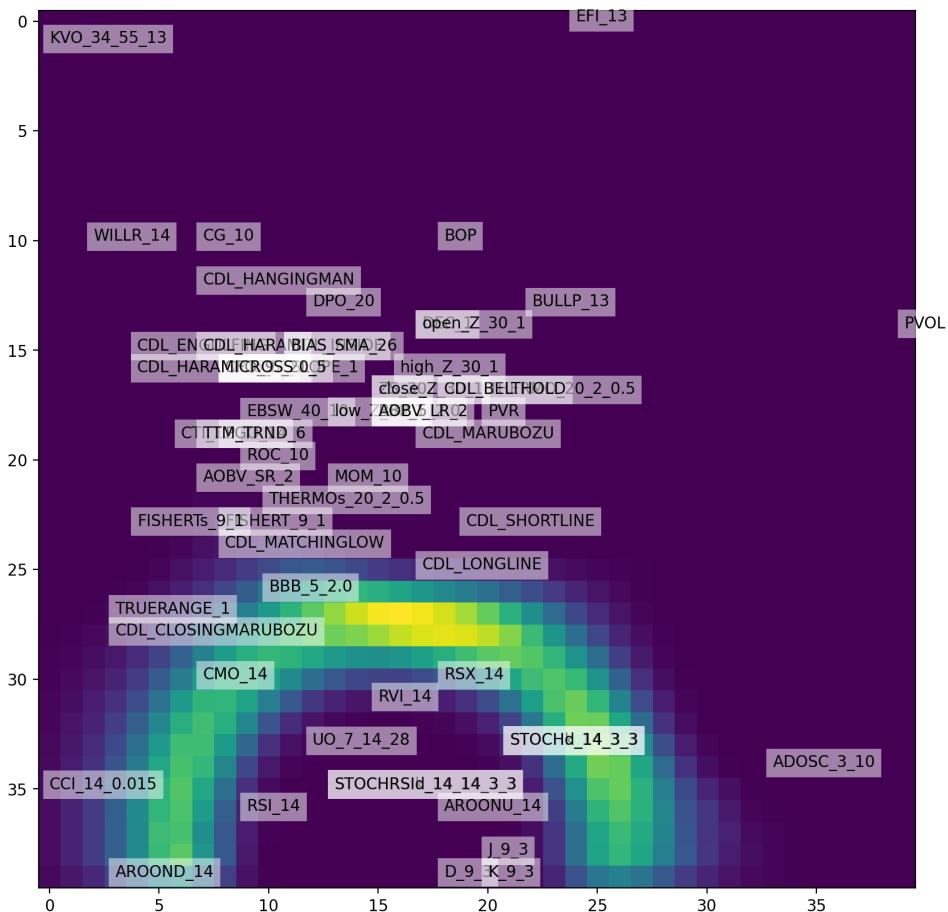
$\theta(t)$  is the neighbourhood function

First of all, the initialised dimension of SOM should be determined. The recommendation dimension about the rule of thumb is based on [Kannan Singaravelu \(2022d\)](#)

$$Q^* \approx 5 \times \sqrt{I}$$

Where  $I$  is the number of features. The data includes 64 features in this part. Therefore,  $Q^*$  is equal to 40. The initialised dimension of SOM is  $40 \times 40$ . The visualisation of SOM is shown by Figure 7. The big purple area contains the majority of the features. Following the SOM selection criteria, the selected features are shown by Table 1.

Figure 7: Self Organising Maps



## 4 Exploratory Data Analysis

Exploratory Data Analysis (EDA) refers to the critical process of performing initial investigations on data so as to discover patterns, spot anomalies, test hypotheses, check assumptions, etc [Kannan Singaravelu \(2022a\)](#). Several important techniques can be used to describe the significant statistical analysis and graphical representations, including correlation matrix, pairwise plot, etc.

Table 1: Features List

Features	
1	KVO_34_55_13
2	EBSW_40_10
3	DPO_20
4	WILLR_14
5	BULLP_13
6	TRUERANGE_1
7	EFI_13
8	ADOSC_3_10
9	BOP
10	PVR

## 4.1 Features Describe

Figure 8 shows the descriptions of all selected features before applying the scaler. The population of the data is 2899. Several important statistics are described by it, including mean, standard deviation, maximum, etc.

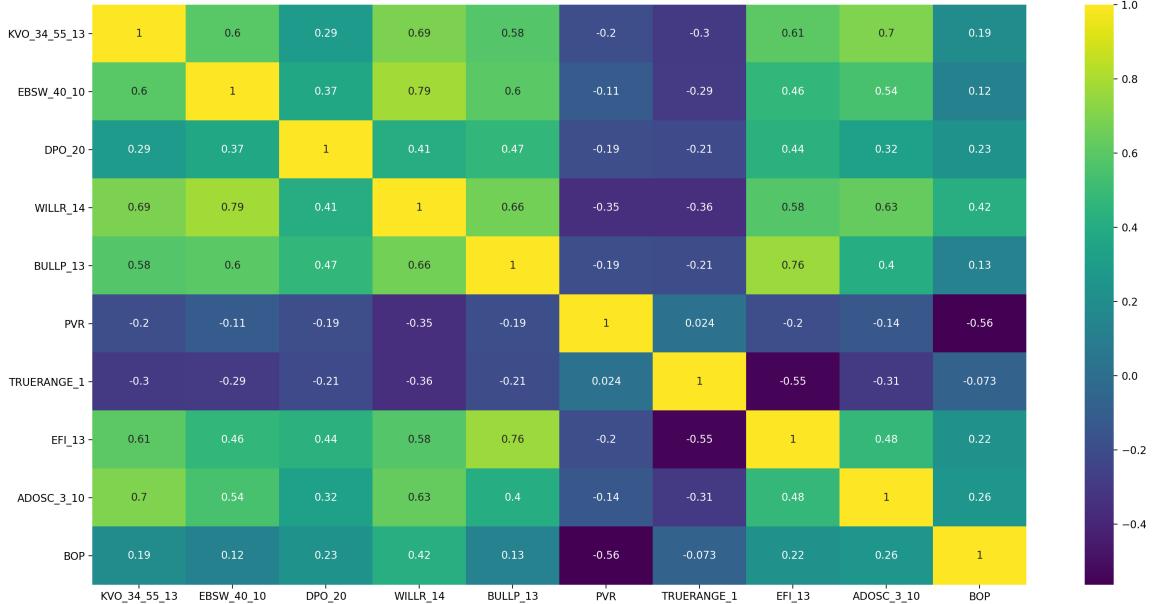
Figure 8: Feature Descriptions

	count	mean	std	min	25%	50%	75%	max
<b>KVO_34_55_13</b>	2899.0	-3.152470e+04	7.442190e+06	-5.814442e+07	-4.040862e+06	4.774819e+05	4.347242e+06	2.995325e+07
<b>EBSW_40_10</b>	2899.0	2.145788e-01	8.124232e-01	-9.999737e-01	-7.566179e-01	6.575463e-01	9.683719e-01	9.999898e-01
<b>DPO_20</b>	2899.0	-1.420945e-01	3.917509e+00	-2.789100e+01	-1.595507e+00	3.150635e-02	1.299499e+00	2.668348e+01
<b>WILLR_14</b>	2899.0	-3.583892e+01	3.061310e+01	-1.000000e+02	-5.989283e+01	-2.692308e+01	-8.134081e+00	-0.000000e+00
<b>BULLP_13</b>	2899.0	1.895417e+00	3.997989e+00	-2.681032e+01	3.935218e-01	1.819845e+00	3.505296e+00	2.179554e+01
<b>PVR</b>	2899.0	2.410486e+00	1.016686e+00	1.000000e+00	2.000000e+00	2.000000e+00	3.000000e+00	4.000000e+00
<b>TRUERANGE_1</b>	2899.0	3.254581e+00	3.039621e+00	4.400024e-01	1.389999e+00	2.190002e+00	3.914993e+00	3.196001e+01
<b>EFI_13</b>	2899.0	-2.309790e+07	1.559325e+08	-2.469719e+09	-4.078148e+07	7.807426e+06	3.476777e+07	4.330140e+08
<b>ADOSC_3_10</b>	2899.0	2.980900e+07	7.917299e+07	-4.047211e+08	-1.544632e+07	3.527280e+07	7.614915e+07	3.343292e+08
<b>BOP</b>	2899.0	5.717779e-02	5.265266e-01	-9.838320e-01	-3.874843e-01	7.758349e-02	5.138909e-01	1.000000e+00

## 4.2 Correlation Matrix

To check the correlation among features and avoid multicollinearity, the correlation matrix was employed shown by Figure 9. Besides, the below code was implemented to check if there are pair variables with absolute correlation exceeding  $\pm 0.8$ , which proves that there are no variables with higher correlations.

Figure 9: Feature Correlation Matrix



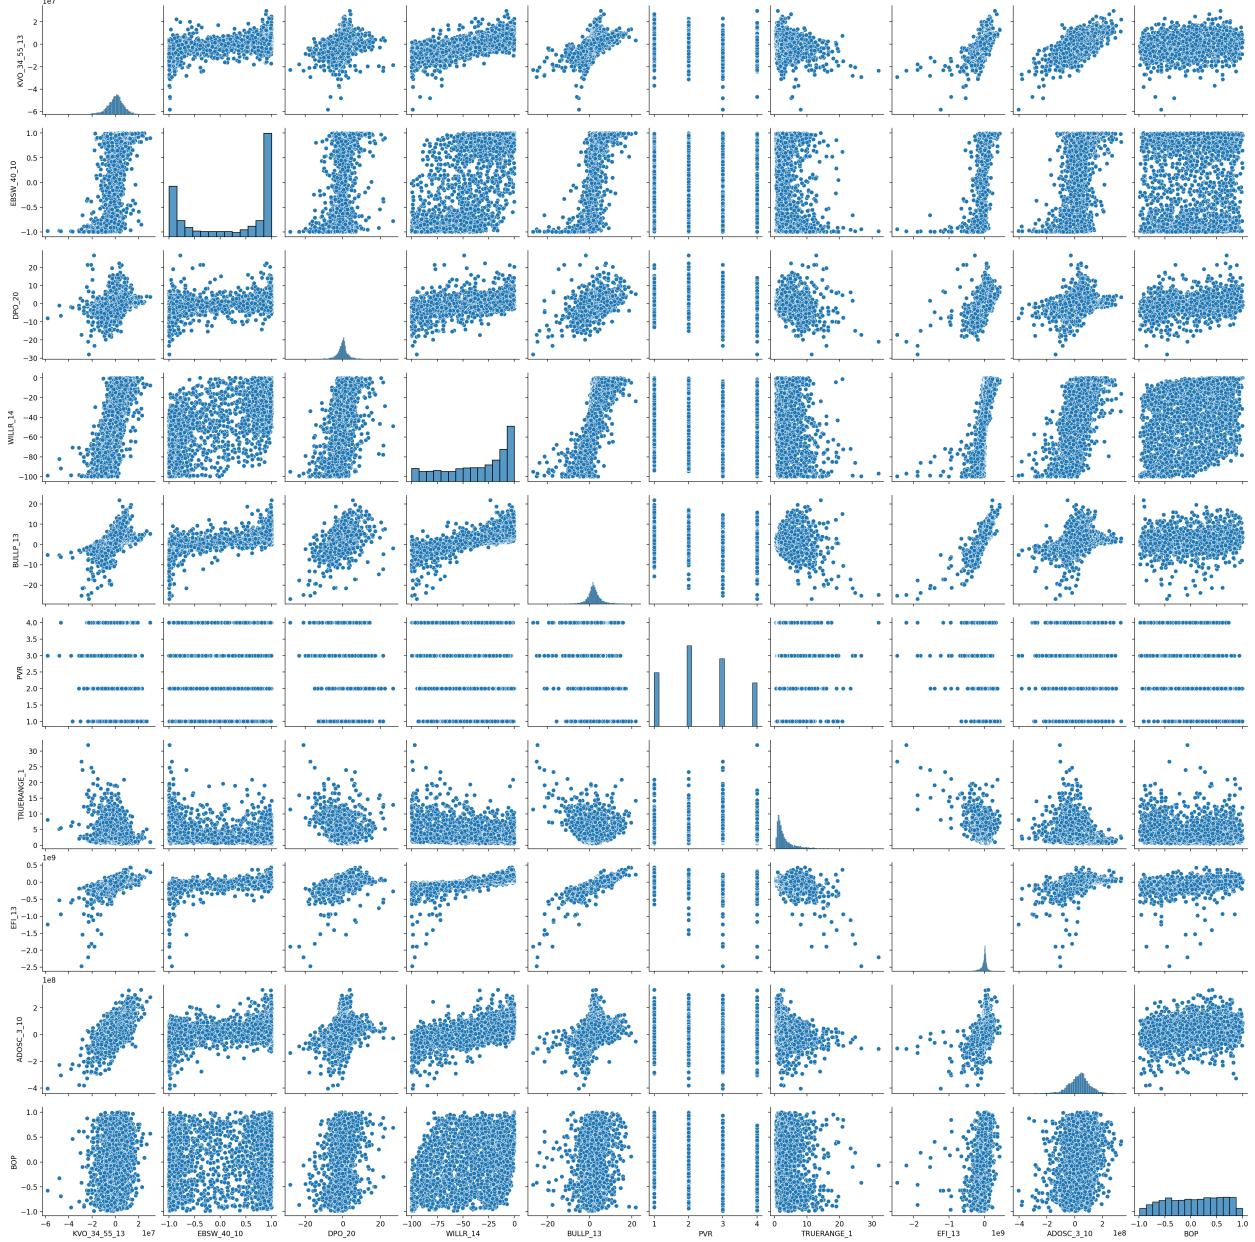
```
# setting correlation maximum threshold
rho_max = 0.80

# detecting pairwise features with abs(corr) > rho_max
idxPairwise = np.where(np.abs(corrFeatMatrix) >= rho_max)
lstPairwise = [
    [corrFeatMatrix.index[x], corrFeatMatrix.columns[y],
     round(corrFeatMatrix.iloc[x, y], 3)]
    for x, y in zip(*idxPairwise) if x != y and x < y]
print(":::>> Pairwise features with correlation = 0.80 :")
print("::::>>> List/list represents [Feat1, Feat2, Corr]:")
```

### 4.3 Pairwise Plot

The pairwise plot is another method to describe the relationship among the features, which calculates different subplots to see if there are differentiable relationships between each pair with respect to the labels. For instance, there is a positive relationship between 'ADOSC\_3\_10' and 'KVO\_34\_55\_13'. Besides, the relationship between 'EBSW\_40\_10' and 'BULLP\_13' is close to logistic regression.

Figure 10: Feature Pairwise Plot



## 5 Experiments And Results

There are some factors to impact the performance of the Deep Neural Network, including the number of hidden layers, the type of selected activation, and relevant parameters. Therefore, three models are employed in this project, Simple LSTM, Stacked-LSTM, and LSTM with Hyperparameter Optimization to evaluate the performance, develop an optimal LSTM model for stock trend projections, and construct the trading strategies. In the first two models, MSE and MAE are the loss functions to optimise the model shown by Equation 1 and Equation 2.

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2 \quad (1)$$

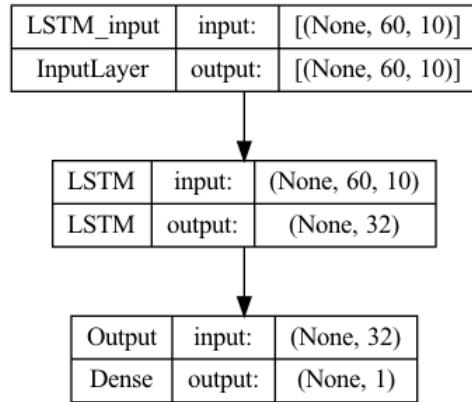
$$\text{MAE} = \frac{\sum_{i=1}^N (y_i - \hat{y}_i)}{N} \quad (2)$$

## 5.1 Simple LSTM

The simple LSTM model structure is shown in Figure 11. The lookback period is 60 days. It only has one hidden layer with the activation function of "relu", which is defined as Equation 3.

$$f(x) = x^+ = \max(0, x) \begin{cases} x, & x > 0 \\ 0, & \text{Otherwise} \end{cases} \quad (3)$$

Figure 11: Simple LSTM Model



Python code for the single LSTM model is as follows

```

# Create a model
def create_model(hu=256, lookback=60, features=numfeat):
    tensorflow.keras.backend.clear_session()
    # instantiate the model
    model = Sequential()
    model.add(LSTM(units=hu, input_shape=(lookback, features),
                   activation = 'relu', return_sequences=False, name='LSTM'))
    model.add(Dense(units=1, name='Output'))
    # specify optimizer separately (preferred method)
    opt = RMSprop(learning_rate=0.001, rho=0.9, epsilon=1e-08)
    # model compilation
    model.compile(optimizer=opt, loss='mse', metrics=['mse'])
    return model

```

The confusion matrix of the simple LSTM model is shown by Figure 12, which displays the detailed classification. And Figure 13 shows the Area under the ROC curve (AUC) is 0.62,

representing a binary classification model's ability to separate positive classes from negative classes. Moving on to the classification report, the precision indicates the proportion of positive identifications (model predicted class 1) which were actually correct. A model which produces no false positives has a precision of 1.0. Recall indicates the proportion of actual positives which were correctly classified. A model which produces no false negatives has a recall of 1.0. The F1 score is the combination of precision and recall. A perfect model achieves an F1 score of 1.0. Support is the number of samples for each class. From this classification report, the accuracy is 0.64, which is just much better than random. The values of the macro avg are the same as the weighted avg, meaning that our results are not damaged by the imbalanced class of the sample.

Figure 12: Simple LSTM Confusion Matrix

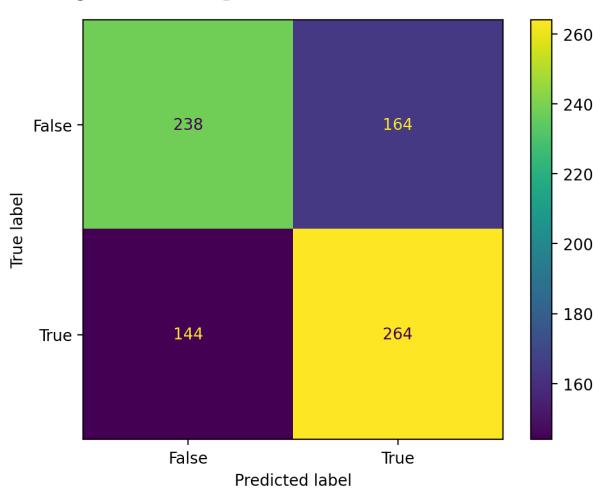
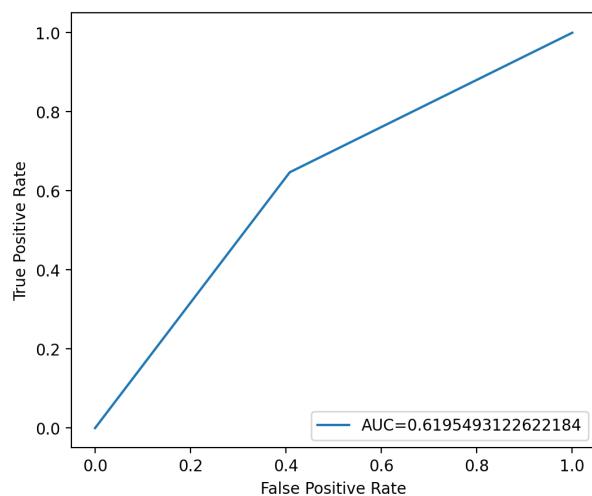


Figure 13: Simple LSTM ROC



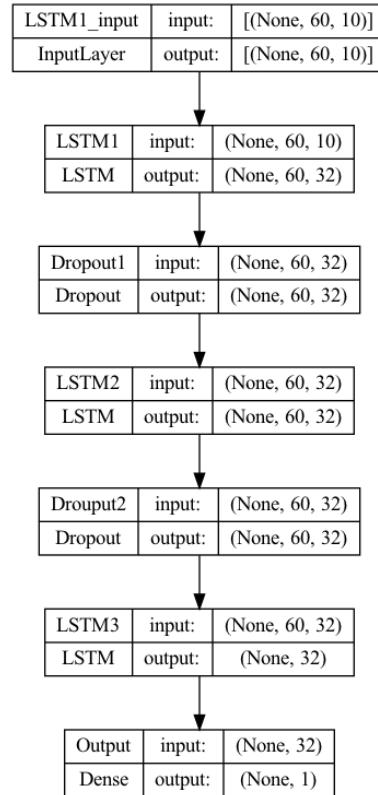
#### # The classification report:

	precision	recall	f1-score	support
0	0.62	0.59	0.61	402
1	0.62	0.65	0.63	408
accuracy			0.62	810
macro avg	0.62	0.62	0.62	810
weighted avg	0.62	0.62	0.62	810

## 5.2 Stacked-LSTM Model

The Stacked-LSTM Model was created by adding the number of hidden layers with 3 hidden layers, relu, gelu, and selu respectively shown by Figure 14.

Figure 14: Stacked-LSTM Model



Python code for the Single LSTM model is as follows

```

# Create a model
def create_model(hu=256, lookback=60, features=numfeat):
    tensorflow.keras.backend.clear_session()
    # instantiate the model
    model = Sequential()
    model.add(LSTM(units=hu, input_shape=(lookback, features),
                   activation = 'relu', return_sequences=True, name='LSTM1'))
    model.add(Dropout(0.4, name='Dropout1'))
    model.add(LSTM(units=hu, activation = 'gelu',
                   return_sequences=True, name='LSTM2'))
    model.add(Dropout(0.4, name='Dropout2'))
    model.add(LSTM(units=hu, activation = 'selu',
                   return_sequences=False, name='LSTM3'))
    model.add(Dense(units=1, activation='sigmoid', name='Output'))
    # specify optimizer separately (preferred method)
    opt = RMSprop(learning_rate=0.001, rho=0.9, epsilon=1e-08)
    # model compilation
    model.compile(optimizer=opt, loss='mse', metrics=['mae'])
    return modelizer=opt, loss='mse', metrics=['mse'])
    return model
  
```

Figure 15 and Figure 16 show the accuracy from the Stacked-LSTM Model has a little improvement since its AUC of 0.63 is more than 0.62 of the Simple LSTM Model by adding the number of hidden layers. Besides, executed time is another key point. Computing speed is another key point. We need to trade off the model accuracy and computing speed.

Figure 15: Stacked LSTM Confusion Matrix

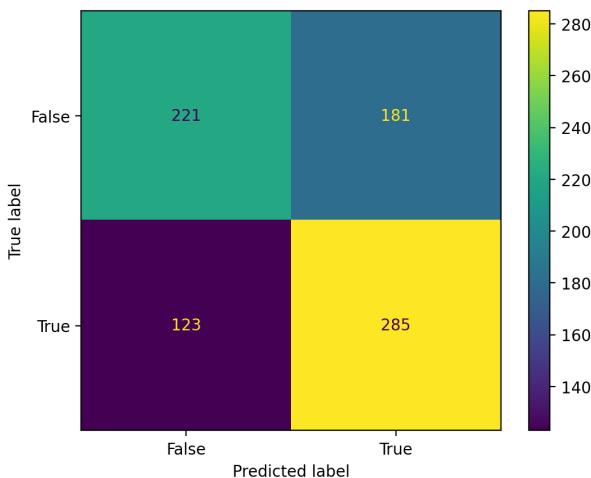
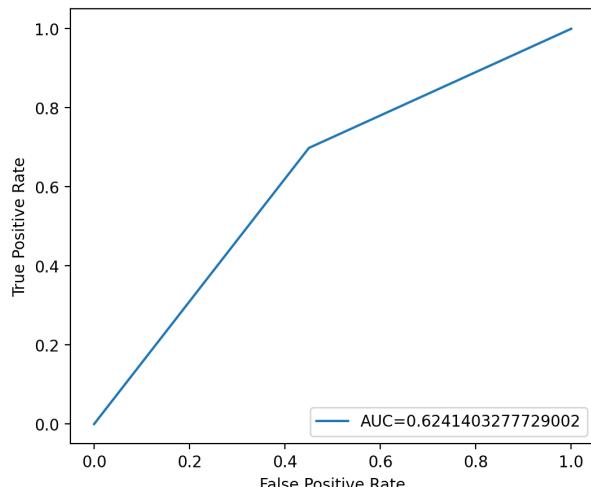


Figure 16: Stacked LSTM ROC



```
# The classification report:
```

	precision	recall	f1-score	support
0	0.64	0.55	0.59	402
1	0.61	0.70	0.65	408
accuracy			0.62	810
macro avg	0.63	0.62	0.62	810
weighted avg	0.63	0.62	0.62	810

### 5.3 Hyperparameter Optimization LSTM

Hyperparameter is a training parameter whose value is used to control the learning process as compared to the values of other parameters (typically weights) that are learned; set before training the model and are not learned by the model during the training process while the model parameters are learned by the model, from the data, during the training process. Hyperparameter Optimization is another the most common method to optimise for better results by choosing hidden layers, number of neurons, choice of activation function, number of epochs, learning rate, batch size and regularisation parameters [Kannan Singaravelu \(2022c\)](#). There are three methods for Hyperparameter Optimization, Random Search, HyperBand, and Bayesian Optimization. In this project, Bayesian Optimization is employed to optimise the Hyperparameter. The best parameters are as follows

```

Trial 15 Complete [00h 00m 56s]
val_accuracy: 0.644444465637207
Best val_accuracy So Far: 0.6567901372909546
Total elapsed time: 00h 09m 31s
INFO:tensorflow:Oracle triggered exit
{'units1': 32, 'units2': 32, 'units3': 32, 'Dropout_rate': 0.2,
'learning_rate': 0.001, 'activation': 'elu'}

```

Figure 17: Hyperparameter Optimization LSTM

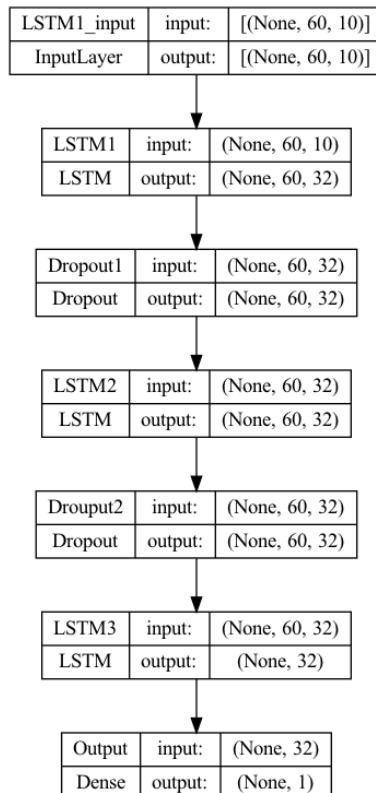


Figure 18 and Figure 19 show the results from the Stacked-LSTM Model with three hidden layers. From Figure 18, the true classification number is larger than the Simple Model. Meanwhile, its AUC is also better equal to 0.64. It proves that Hyperparameter Optimization can improve the accuracy of the LSTM model.

Figure 18: Optimised LSTM Confusion Matrix

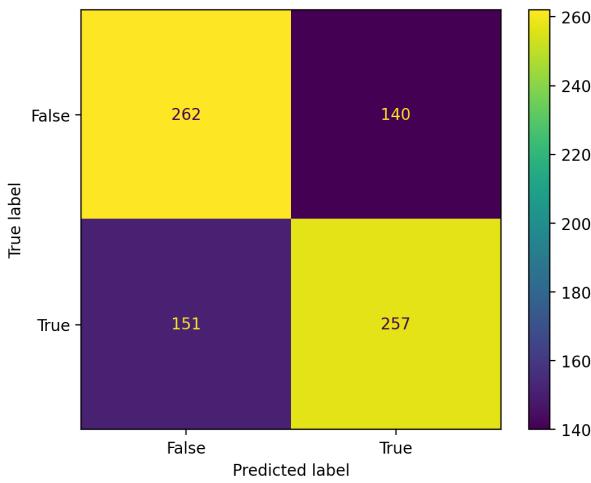
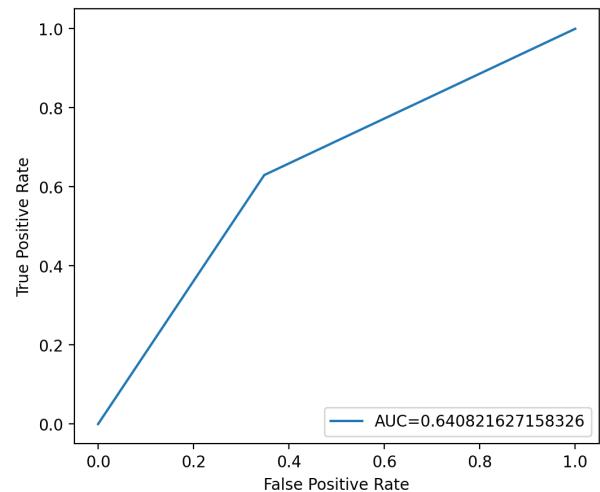


Figure 19: Optimised LSTM ROC



```
# The classification report:
```

	precision	recall	f1-score	support
0	0.63	0.65	0.64	402
1	0.65	0.63	0.64	408
accuracy			0.64	810
macro avg	0.64	0.64	0.64	810
weighted avg	0.64	0.64	0.64	810

## 6 Trading Strategies

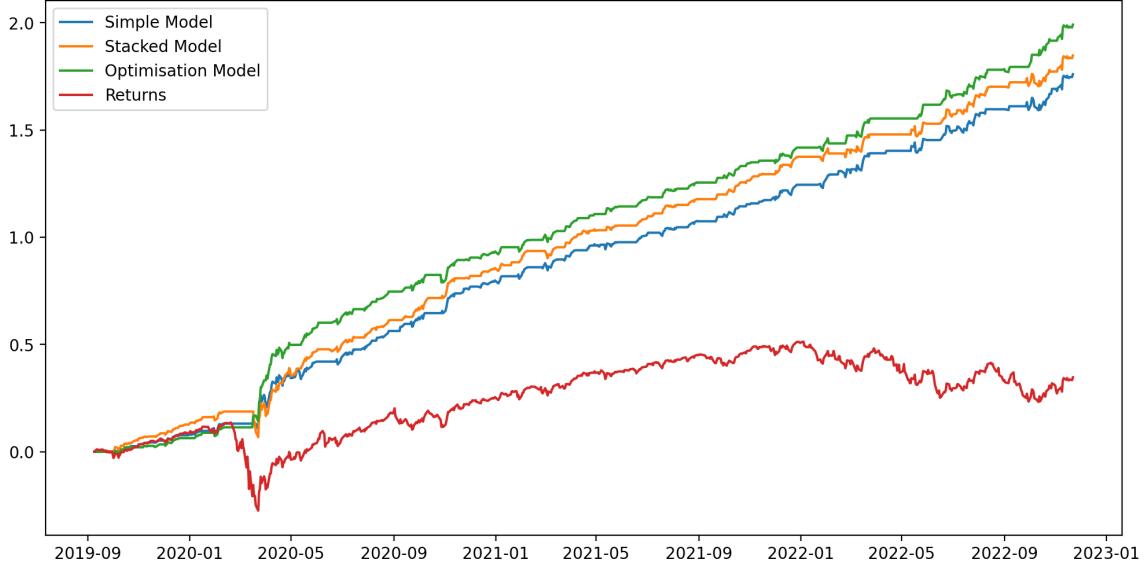
### 6.1 Long-Only Strategy

For the prediction for each stock on the next day, if the prediction is positive, we buy the stock at the open price and sell the stock at a close price on the same day. If the prediction is negative, no action is taken. Figure 20 shows the performance of applying the model projections to the strategies during the test periods. The cumulative returns are shown by Table 2. Moreover, the LSTM with Hyperparameter Optimisation has the highest cumulative returns. Meanwhile, its accuracy is highest in these three models.

Table 2: Long-Only Returns

Models	Cumulative Returns
Single LSTM	176.01%
Stacked LSTM	184.74%
Optimisation LSTM	199.13%
Holding	34.67%

Figure 20: Long-Only Strategy Return



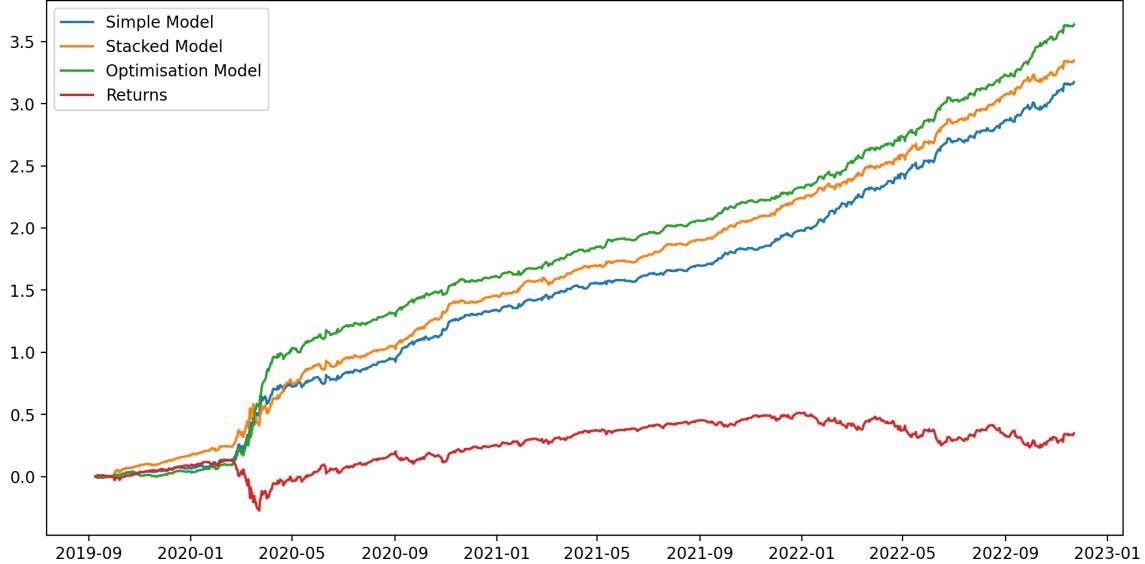
## 6.2 Long-Short Strategy

For the prediction for each stock on the next day, if the prediction is positive, we buy the stock at the open price and sell the stock at a close price on the same day. If the prediction is negative, we short-sell the stock at the open price and close out the short-sell at the close price. Figure 21 shows the cumulative returns of three strategies from three LSTM models. The results are the same as the Long-Only strategy. However, the cumulative returns of the Long-Short strategy outstrip the Long-Only strategy. Table 3 shows the detailed results.

Table 3: Long-Short Returns

Models	Cumulative Returns
Single LSTM	317.34%
Stacked LSTM	334.80%
Optimisation LSTM	363.58%
Holding	34.67%

Figure 21: Long-Short Strategy Return



## 7 Conclusion

This project establishes a forecasting framework to predict the direction (Up/Down) of stocks. I leveraged the combinations of price, volumes and corporate statistics as input data by Technical Analysis Library. Three models are proposed, developed, trained, and tested: Single LSTM, Stacked-LSTM, and LSTM with optimised hyperparameter. And two trading strategies are created according to the model predictions. The Optimised hyperparameter LSTM shows more superior results over other models due its its ability to search the best parameters of the LSTM models. Hence the Optimised hyperparameter LSTM is more able to capture the information in the time series and more suitable in predicting financial time series. Our superior trading return from the Optimised hyperparameter LSTM further validates our experimental result. Moreover, we have shown that despite the more complicated model structure of stacked-LSTM over single LSTM model, the stacked-LSTM has better model performance over the single LSTM model. One drawback of the Stacked-LSTM might cause overfitting, which can be avoid by Ridge or LASSO regression.

## Table of Contents

- [1 Import Libraries](#)
- [2 Loading the Historical Data](#)
- [3 Features Selection](#)
  - [3.1 Borutapay](#)
  - [3.2 K-Means clustering](#)
    - [3.2.1 Elbow Plot](#)
    - [3.2.2 Silhouette Coefficient](#)
  - [3.3 SOM](#)
- [4 Exploratory Data Analysis](#)
  - [4.1 Features Describe](#)
  - [4.2 Correlation Matrix](#)
  - [4.3 Pair Plot](#)
  - [4.4 XGBoosting](#)
  - [4.5 Spliting](#)
- [5 LSTM Model](#)
  - [5.1 Model 1 LSTM Model](#)
  - [5.2 Model 2 Stacked-LSTM Model](#)
- [6 Hyperparameter Optimization](#)
  - [6.1 Approach 1 : Random Search](#)
  - [6.2 Approach 2 : HyperBand](#)
  - [6.3 Approach 3 : BayesianOptimization](#)
- [7 Optimal LSTM](#)
- [8 Trading Strategy](#)
  - [8.1 Long-Only Strategy](#)
  - [8.2 Long-Short Strategy](#)

# Deep Learning For Financial Time Series

## 1 Import Libraries

```
In [1]: 1 # Base Libraries
2 import numpy as np
3 import pandas as pd
4 import datetime
5 from datetime import datetime
6 import copy
7
8 import yfinance as yf
9 import cufflinks as cf
10 cf.set_config_file(offline=True, dimensions=(1000,600))
11
12 # Visualization
13 import matplotlib
14 import matplotlib.pyplot as plt
15 # # Plot settings
16 # plt.style.use('seaborn-white')
17 # matplotlib.rcParams['figure.figsize'] = [12.0, 6.0]
18 # matplotlib.rcParams['font.size'] = 10
19 # matplotlib.rcParams['lines.linewidth'] = 2.0
20 # matplotlib.rcParams['grid.color'] = 'white'
21 import seaborn as sns
22 # plt.style.use('seaborn')
23 %matplotlib inline
24 %config InlineBackend.figure_format = 'retina'
25
26 import ta
27 from ta import add_all_ta_features
28 from ta.utils import dropna
29
30 # Preprocessing
31 from sklearn.pipeline import Pipeline
32 from sklearn.preprocessing import MinMaxScaler, RobustScaler, StandardScaler, Normalizer
33 from sklearn.model_selection import train_test_split, TimeSeriesSplit, GridSearchCV
34 from sklearn.model_selection import RandomizedSearchCV, cross_val_score
35 from sklearn.metrics import mean_squared_error, r2_score
36 from sklearn.metrics import roc_curve, auc
37 from sklearn import metrics
38
39 # Import sklearn modules
40 from sklearn.model_selection import train_test_split
41 from sklearn.cluster import KMeans
42 from sklearn.pipeline import Pipeline
43
44 # SOM & Sklearn library
45 from minisom import MiniSom
46
47 # XGBoost Classifier
48 from xgboost import XGBClassifier, plot_importance, to_graphviz
49
50 # import shap
51
52 import datetime
53 from pathlib import Path
54
55 results_path = Path('results', 'lstm_time_series')
56 if not results_path.exists():
57     results_path.mkdir(parents=True)
58
59 # imports
60 from src.config import *
61 pd.set_option('display.max_columns', 30)
62 pd.set_option('display.max_rows', 1000)
63
64 # import tensorflow and keras
65 import tensorflow
66 # tensorflow modules
67 from tensorflow.keras.models import Sequential
68 from tensorflow.keras.layers import Dense, Dropout, Flatten, LSTM
69 from tensorflow.keras.optimizers import Adam, RMSprop
70 from tensorflow.keras.utils import plot_model
71 from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, TensorBoard
72
73 # Load the TensorBoard notebook extension
74 %load_ext tensorboard
75
76 # Ignore warnings
77 import warnings
78 warnings.filterwarnings('ignore')
```

executed in 8.70s, finished 20:07:44 2023-01-23



```

In [2]: 1 # def set_seeds(seed=2023):
2 #     random.seed(seed)
3 #     np.random.seed(seed)
4 #     tensorflow.set_random_seed(seed)
5
6 # Plot a confusion matrix.
7 # cm is the confusion matrix, names are the names of the classes.
8 def plot_confusion_matrix(cm, names, title='Confusion matrix', cmap=plt.cm.Blues):
9     plt.imshow(cm, interpolation='nearest', cmap=cmap)
10    plt.title(title)
11    plt.colorbar()
12    tick_marks = np.arange(len(names))
13    plt.xticks(tick_marks, names, rotation=45)
14    plt.yticks(tick_marks, names)
15    plt.tight_layout()
16    plt.ylabel('True label')
17    plt.xlabel('Predicted label')
18
19 # Plot an ROC. pred - the predictions, y - the expected output.
20 def plot_roc(pred,y):
21     fpr, tpr, _ = roc_curve(y, pred)
22     roc_auc = auc(fpr, tpr)
23
24     plt.figure()
25     plt.plot (fpr, tpr, label='ROC curve (area = %0.2f) ' % roc_auc)
26     plt.plot ([0, 1], [0, 1], 'k--')
27     plt.xlim([0.0, 1.0])
28     plt.ylim([0.0, 1.05])
29     plt.xlabel('False Positive Rate')
30     plt.ylabel('True Positive Rate')
31     plt.title('Receiver Operating Characteristic (ROC)')
32     plt.legend(loc="lower right")
33     plt.show()

```

executed in 5ms, finished 20:07:44 2023-01-23

## 2 Loading the Historical Data

```

In [3]: 1 set_seeds(2023)
executed in 3ms, finished 20:07:44 2023-01-23

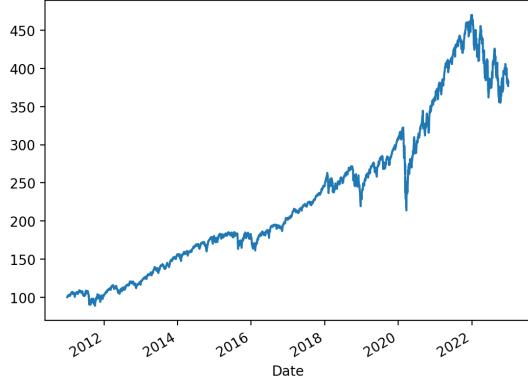
In [4]: 1 ticker = 'SPDR S&P 500 ETF Trust'
executed in 2ms, finished 20:07:44 2023-01-23

In [5]: 1 row_data = yf.download('SPY', start='2010-12-31',end='2022-12-31')
executed in 575ms, finished 20:07:45 2023-01-23
[******100%*****] 1 of 1 completed

In [105]: 1 # row_data['Adj Close'].plot(title= ticker + ' Adjusted Close')
2 row_data['Adj Close'].plot()
executed in 249ms, finished 20:30:20 2023-01-23

```

Out[105]: <AxesSubplot:xlabel='Date'>



```

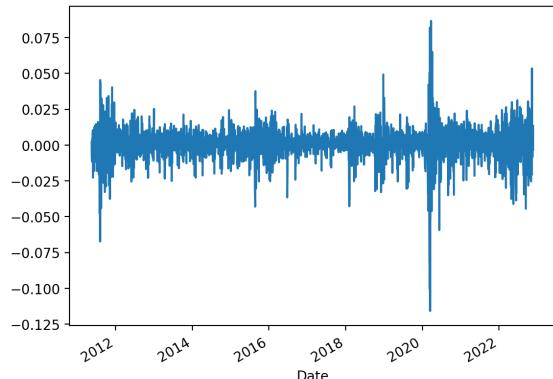
In [7]: 1 threshold = 0.09/100 # Expense Ratio
executed in 3ms, finished 20:07:45 2023-01-23

In [8]: 1 row_data.ta.strategy('All')
executed in 3.19s, finished 20:07:48 2023-01-23
131it [00:03, 43.51it/s]

In [9]: 1 stock_data = row_data.copy()
executed in 6ms, finished 20:07:48 2023-01-23

In [106]: 1 stock_data['Returns'] = np.log(row_data['Adj Close']).diff().fillna(0)
2 # stock_data['Returns'].plot(title= ticker + ' Daily Returns')
3 stock_data['Returns'].plot()
4 stock_data['Labels'] = stock_data['Returns'].apply(lambda x: 0 if x<=threshold else 1)
executed in 279ms, finished 20:30:39 2023-01-23

```



```
In [11]: 1 stock_data.query('Labels == 1').shape[0]
executed in 21ms, finished 20:07:49 2023-01-23
Out[11]: 1459
```

```
In [12]: 1 label_list = ['0', '1']
2 count0 = stock_data.query('Labels == 0').shape[0]
3 count1 = stock_data.query('Labels == 1').shape[0]
4 height = [count0, count1]
5 plt.bar(label_list, height=height)
executed in 155ms, finished 20:07:49 2023-01-23
Out[12]: <BarContainer object of 2 artists>
```

```
In [13]: 1 stock_data.drop(['HILOL_13_21', 'HILOS_13_21', 'PSARL_0_02_0_2',
2 'PSARS_0_02_0_2', 'PSARaf_0_02_0_2', 'QOEL_14_5_4.236',
3 'QQEs_14_5_4.236', 'SUPERTl_7_3.0', 'SUPERTs_7_3.0',
4 'Open', 'High', 'Low', 'Close', 'Adj Close'],
5 axis=1, inplace=True)
executed in 6ms, finished 20:07:49 2023-01-23
```

```
In [14]: 1 stock_data.describe().T
executed in 496ms, finished 20:07:50 2023-01-23
Out[14]:
```

	count	mean	std	min	25%	50%	75%	max
Volume	3021.0	1.106428e+08	6.434431e+07	2.027000e+07	6.805420e+07	9.365700e+07	1.344574e+08	7.178287e+08
ABER_ZG_5_15	3017.0	2.503349e+02	9.518406e+01	1.130060e+02	1.807540e+02	2.263233e+02	2.998340e+02	4.767867e+02
ABER_SG_5_15	3006.0	2.539932e+02	9.661298e+01	1.163644e+02	1.838875e+02	2.280344e+02	3.033147e+02	4.821783e+02
ABER_XG_5_15	3006.0	2.475736e+02	9.355612e+01	1.096476e+02	1.809428e+02	2.248068e+02	2.968555e+02	4.713950e+02
ABER_ATR_5_15	3006.0	3.209786e+00	2.331560e+00	9.770435e-01	1.592276e+00	2.189555e+00	3.996600e-00	1.498682e+01
ACCBBL_20	3002.0	2.431010e+02	9.147485e+01	1.077167e+02	1.773469e+02	2.220095e+02	2.913095e+02	4.572807e+02
ACCBM_20	3002.0	2.503653e+02	9.489745e+01	1.163625e+02	1.810781e+02	2.259125e+02	2.980766e+02	4.699965e+02
ACCBU_20	3002.0	2.575436e+02	9.849327e+01	1.235737e+02	1.848288e+02	2.292262e+02	3.050152e+02	4.832534e+02
AD	3021.0	1.639830e+10	7.824943e+09	5.067631e+07	1.064866e+10	1.780189e+10	2.331822e+10	2.700806e+10
ADOSC_3_10	3012.0	3.083129e+07	7.932924e+07	-4.047211e+08	-1.470026e+07	3.563843e+07	7.706414e+07	3.343292e+08
ADX_14	2994.0	2.180187e+01	7.506001e+00	8.593231e+00	1.595191e+01	2.066805e+01	2.651445e+01	5.300331e+01

```
In [15]: 1 stock_data.shape
executed in 3ms, finished 20:07:50 2023-01-23
Out[15]: (3021, 272)
```

```
In [16]: 1 stock_data.isna().sum()
executed in 10ms, finished 20:07:50 2023-01-23
Out[16]:
```

Volume	0
ABER_ZG_5_15	4
ABER_SG_5_15	15
ABER_XG_5_15	15
ABER_ATR_5_15	15
ACCBBL_20	19
ACCBM_20	19
ACCBU_20	19
AD	0
ADOSC_3_10	9
ADX_14	27
DMP_14	13
DMN_14	13
ALMA_10_6_0_0.85	10
AMATE_LR_8_21_2	0
AMATE_SR_8_21_2	0
AO_5_34	33
OBV	0
OBV_min_2	1
***_***_***	*

```
In [17]: 1 # drop nan values
2 stock_data = stock_data.dropna(axis=0)
executed in 8ms, finished 20:07:50 2023-01-23
```

```
In [18]: 1 stock_data.isna().sum()
executed in 9ms, finished 20:07:50 2023-01-23
```

```
Out[18]: Volume          0
ABER_ZG_5_15      0
ABER_SG_5_15      0
ABER_XG_5_15      0
ABER_ATR_5_15      0
ACCBL_20          0
ACCBM_20          0
ACCBU_20          0
AD              0
ADOSC_3_10        0
ADX_14            0
DMP_14            0
DMN_14            0
ALMA_10_6.0_0.85  0
AMATE_LR_8_21_2    0
AMATE_SR_8_21_2    0
AO_5_34            0
OBV              0
OBV_min_2         0
----
```

## 3 Features Selection

```
In [19]: 1 y = stock_data.iloc[:, :-1]
2 X = stock_data.iloc[:, :-2]
executed in 3ms, finished 20:07:50 2023-01-23
```

```
In [20]: 1 # class frequency
2 c = stock_data['Returns'].value_counts()
3
4 # class weight function
5 def cwts(stock_data):
6     c0, c1 = np.bincount(stock_data['Labels'])
7     w0=(1/c0)*(len(stock_data['Returns']))/2
8     w1=(1/c1)*(len(stock_data['Returns']))/2
9     return {0: w0, 1: w1}
10
11 # check class weights
12 class_weight = cwts(stock_data)
13 class_weight
executed in 7ms, finished 20:07:50 2023-01-23
```

```
Out[20]: {0: 0.966977985323549, 1: 1.0353571428571429}
```

### 3.1 Borutapy

```
In [21]: 1 # define random forest classifier
2 forest = RandomForestClassifier(n_jobs=-1,
3                                 class_weight=cwts(stock_data),
4                                 random_state=42,
5                                 max_depth=5)
6
7 # train the model
8 forest.fit(X, y)
executed in 277ms, finished 20:07:50 2023-01-23
```

```
Out[21]: RandomForestClassifier(class_weight={0: 0.966977985323549,
                                              1: 1.0353571428571429},
                                 max_depth=5, n_jobs=-1, random_state=42)
```

```
In [22]: 1 # define Boruta feature selection method
2 feat_selector = BorutaPy(forest, n_estimators='auto', verbose=2, random_state=0)
3
4 # find all relevant features
5 # takes input in array format not as dataframe
6 X_ary = np.array(X)
7 y_ary = np.array(y)
8 feat_selector.fit(X_ary, y_ary)
9
10 # check selected features
11 print(feat_selector.support_)
12
13 # check ranking of features
14 print(feat_selector.ranking_)
15
16 # call transform() on X to filter it down to selected features
17 X_filtered = feat_selector.transform(X_ary)
executed in 1m 5.21s, finished 20:08:56 2023-01-23
```

```
Iteration: 1 / 100
Confirmed: 0
Tentative: 270
Rejected: 0
Iteration: 2 / 100
Confirmed: 0
Tentative: 270
Rejected: 0
Iteration: 3 / 100
Confirmed: 0
Tentative: 270
Rejected: 0
Iteration: 4 / 100
Confirmed: 0
Tentative: 270
Rejected: 0
Iteration: 5 / 100
Confirmed: 0
Tentative: 270
- . , ^
```

```
In [23]: 1 # zip my names, ranks, and decisions in a single iterable
2 feature_names = X.columns
3 feature_ranks = list(zip(feature_names,
4                         feat_selector.ranking_,
5                         feat_selector.support_))
6
7 # iterate through and print out the results
8 for feat in feature_ranks:
9     print(f'Feature: {feat[0]}: Rank: {feat[1]}: Keep: {feat[2]}')
executed in 12ms, finished 20:08:56 2023-01-23
```

Feature	Rank	Keep
Volume	3	False
ABER_ZG_5_15	108	False
ABER_SC_5_15	163	False
ABER_XG_5_15	97	False
ABER_ATR_5_15	49	False
ACCBEL_20	114	False
ACCBM_20	104	False
ACCBU_20	129	False
AD	89	False
ADOSC_3_10	1	True
ADX_14	65	False
DMP_14	26	False
DNN_14	48	False
ALMA_10_6.0_0.85	134	False
AMATE_LR_8_21_2	12	False
AMATE_SR_8_21_2	7	False
AO_5_34	17	False
OBV	82	False
OBV_min_2	86	False
- . . . ^	- . . . ^	- . . . ^

```
In [24]: 1 selected_rf_features = pd.DataFrame({'Feature':feature_names,
2                                         'Ranking':feat_selector.ranking_})
3
4 # selected_rf_features#.sort_values(by='Ranking')
5
6 selected_rf_features[selected_rf_features['Ranking']==1]
executed in 16ms, finished 20:08:56 2023-01-23
```

Out[24]:

	Feature	Ranking
9	ADOSC_3_10	1
22	AOBV_LR_2	1
23	AOBV_SR_2	1
25	AROOND_14	1
26	AROONU_14	1
32	BBB_5_2.0	1
33	BBP_5_2.0	1
34	BIAS_SMA_26	1
35	BOP	1
38	CCI_14_0.015	1
48	CDL_BELTHOLD	1

```
In [25]: 1 selected = list(selected_rf_features[selected_rf_features['Ranking']==1]['Feature'])
executed in 3ms, finished 20:08:56 2023-01-23
```

```
In [26]: 1 X_selected = stock_data[selected]
executed in 4ms, finished 20:08:56 2023-01-23
```

```
In [27]: 1 X_selected.shape # The number of feature is 64
executed in 6ms, finished 20:08:56 2023-01-23
```

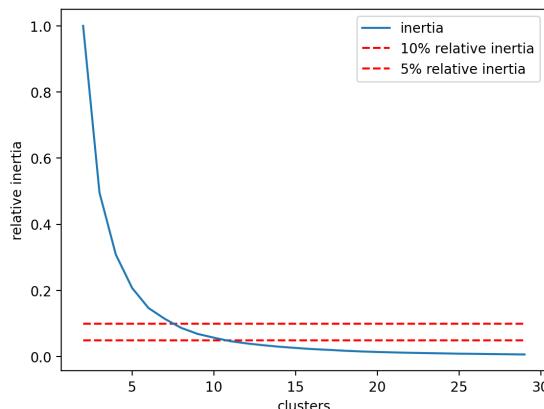
Out[27]: (2899, 64)

## 3.2 K-Means clustering

```
In [28]: 1 X = X_selected
executed in 3ms, finished 20:08:56 2023-01-23
```

### 3.2.1 Elbow Plot

```
In [29]: 1 n_clusters = range(2, 30)
2 inertia = []
3
4 for n in n_clusters:
5     kmeans = KMeans(n_clusters=n)
6     kmeans.fit(X)
7     inertia.append(kmeans.inertia_)
8
9 plt.plot(n_clusters, np.divide(inertia,inertia[0]))
10 plt.hlines(0.1, n_clusters[0], n_clusters[-1], 'r', linestyles='dashed')
11 plt.hlines(0.05, n_clusters[0], n_clusters[-1], 'r', linestyles='dashed')
12 plt.xlabel('clusters')
13 plt.ylabel('relative inertia')
14 plt.legend(['inertia', '10% relative inertia', '5% relative inertia']);
executed in 8.31s, finished 20:09:05 2023-01-23
```



## 3.2.2 Silhouette Coefficient

```
In [30]: 1 from sklearn.cluster import KMeans
2 from sklearn.metrics import silhouette_score
executed in 2ms, finished 20:09:05 2023-01-23
```

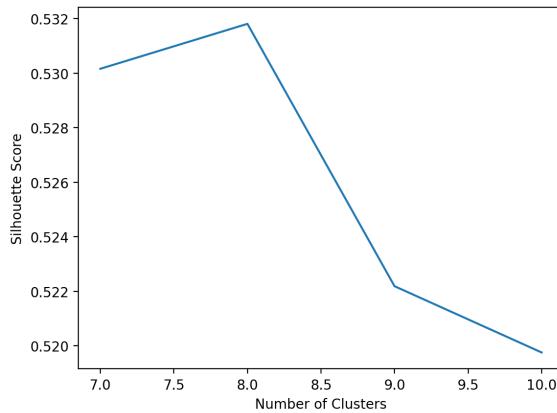
```
In [31]: 1 set_seeds(2023)
2 Silhouette_list = []
3 for i in [7, 8, 9, 10]:
4     KMean_test = KMeans(n_clusters=i)
5     KMean_test.fit(X)
6     label_test = KMean_test.predict(X)
7     score = silhouette_score(X, label_test)
8     Silhouette_list.append(score)
9 #     print(f'Silhouette Score: {silhouette_score(X, label_test)}')
executed in 5.12s, finished 20:09:10 2023-01-23
```

```
In [32]: 1 Silhouette_list
executed in 4ms, finished 20:09:10 2023-01-23
```

```
Out[32]: [0.5301639574706409, 0.5318122369209507, 0.522190364323992, 0.5197598976215557]
```

```
In [33]: 1 Silhouette_label = [7, 8, 9, 10]
2 plt.plot(Silhouette_label, Silhouette_list)
3 plt.xlabel('Number of Clusters')
4 plt.ylabel('Silhouette Score')
executed in 222ms, finished 20:09:10 2023-01-23
```

```
Out[33]: Text(0, 0.5, 'Silhouette Score')
```



```
In [34]: 1 KMean_test = KMeans(n_clusters=7)
2 KMean_test.fit(X)
3 label_test = KMean_test.predict(X)
executed in 210ms, finished 20:09:11 2023-01-23
```

```
In [35]: 1 print(f'Silhouette Score(n=7): {silhouette_score(X, label_test)}')
executed in 138ms, finished 20:09:11 2023-01-23
Silhouette Score(n=7): 0.5310282741634225
```

```
In [36]: 1 # Normalize and fit the model
2 pipe = Pipeline([("normalization", MinMaxScaler()), ("cluster", KMeans(n_clusters=7))])
executed in 3ms, finished 20:09:11 2023-01-23
```

```
In [37]: 1 # Fit Model
2 pipe.fit(X.T)
3
4 # Assign Label
5 labels = pipe.predict(X.T)
executed in 111ms, finished 20:09:11 2023-01-23
```

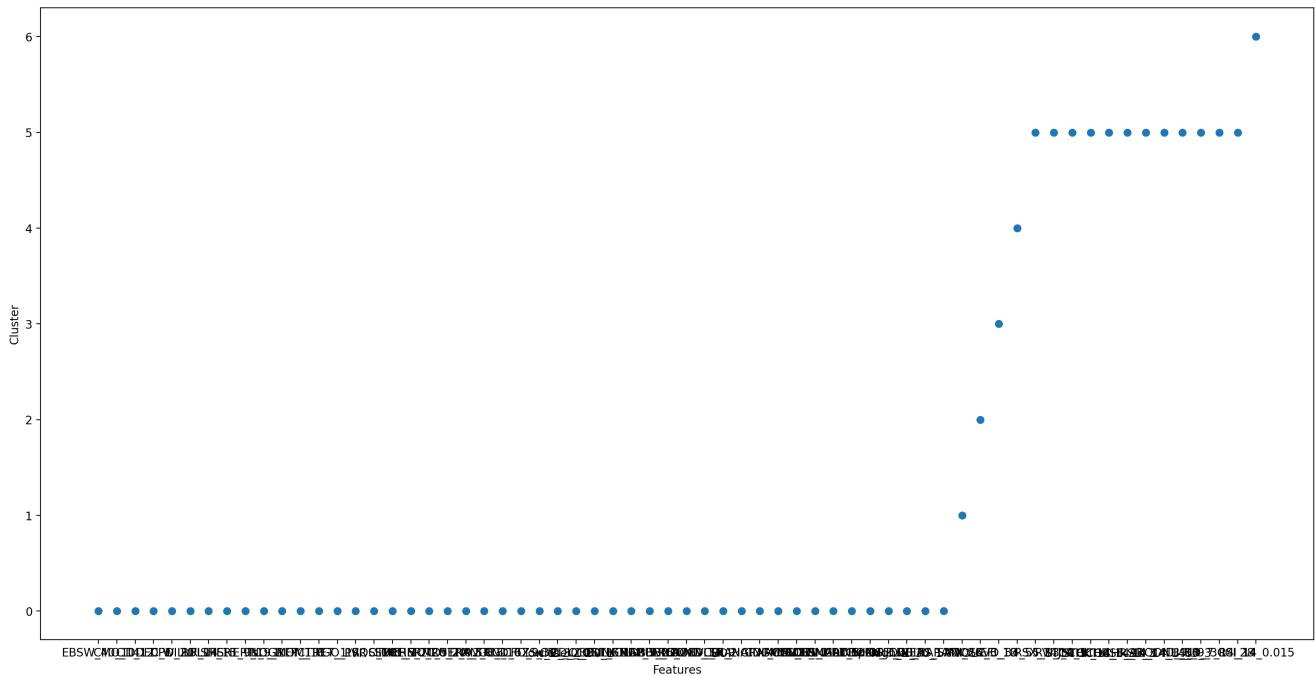
```
In [38]: 1 Features = X.columns
2 # Create dataframe to hold cluster details
3 df = pd.DataFrame({'Cluster': labels,
4                     'Features': Features,
5                     'X_mean': X.mean(axis=0)})
6 df.sort_values(by=['Cluster'], axis=0)
7
8 df = df.reset_index(drop=True)
9 df
10
11 # 'ATR': weekly_atr.mean(axis=1)
executed in 17ms, finished 20:09:11 2023-01-23
```

```
Out[38]:
```

Cluster	Features	X_mean
0	EBSW_40_10	2.145788e-01
1	CMO_14	1.037855e+01
2	CTI_12	2.068876e-01
3	DEC_1	4.522249e-01
4	DPO_20	-1.420945e-01
5	WILLR_14	-3.583892e+01
6	BULLP_13	1.895417e+00
7	FISHERT_9_1	7.843556e-01
8	FISHERTs_9_1	7.832866e-01
9	INC_1	5.443256e-01
10	LOGRET_1	3.762303e-04

```
In [39]: 1 plt.figure(figsize=(20,10))
2 plt.scatter(df.Features, df.Cluster)
3 plt.xlabel('Features')
4 plt.ylabel('Cluster')
5 # plt.title('SPY Features Clustering');
executed in 1.37s, finished 20/09/2023 03-21
```

```
Out[39]: Text(0, 0.5, 'Cluster')
```



```
In [40]: 1 pipe.get_params()  
executed in 7ms, finished 20:09:13 2
```

**Out[40]:** {'memory': None,

```
'steps': [('normalization', MinMaxScaler()),  
          ('cluster', KMeans(n_clusters=7))],  
         'verbose': False,  
         'normalization': MinMaxScaler(),  
         'cluster': KMeans(n_clusters=7),  
         'normalization_clip': False,  
         'normalization_copy': True,  
         'normalization_feature_range': (0, 1),  
         'cluster_algorithm': 'auto',  
         'cluster_copy_x': True,  
         'cluster_init': 'k-means++',  
         'cluster_max_iter': 300,  
         'cluster_n_clusters': 7,  
         'cluster_n_init': 10,  
         'cluster_n_jobs': 'deprecated',  
         'cluster_precompute_distances': 'deprecated',  
         'cluster_random_state': None,  
         'cluster_tol': 0.0001,  
         'cluster_verbose': 0}
```

```
In [41]: 1 pipe['cluster'].cluster_centers_  
executed in 5ms, finished 20:09:13 2023-01-23
```

**Out[41]:** array([ 3.78908751e-04, 1.22

```

    2.70220940e-09, 1.184954592e-09, 5.50018479e-10], ...
    [1.00000000e+00, 1.00000000e+00, 1.00000000e+00, ...
     1.00000000e+00, 1.00000000e+00, 1.00000000e+00], ...
    [5.222050832e-03, 3.12714868e-03, 3.72962350e-03, ...
     8.14431544e-04, 7.10813081e-03, 2.11729830e-03], ...
    ...,
    [0.00000000e+00, 8.32559799e-04, 2.26207208e-03, ...
     1.53280712e-04, 2.01304017e-04, 2.08037790e-04], ...
    [3.78911025e-04, 1.22604386e-03, 2.82096424e-03, ...
     4.68488946e-09, 4.55246113e-09, 3.26746973e-09], ...
    [3.78908452e-04, 1.22603983e-03, 2.82095370e-03, ...
     4.54573273e-09, 3.6413086e-09, 3.27374184e-09]]
```

```
In [42]: 1 pipe['cluster'].cluster centers .shape
```

executed in 4ms, finished 20:09:13 2023-01-23

**Out[42]:** (7, 2899)

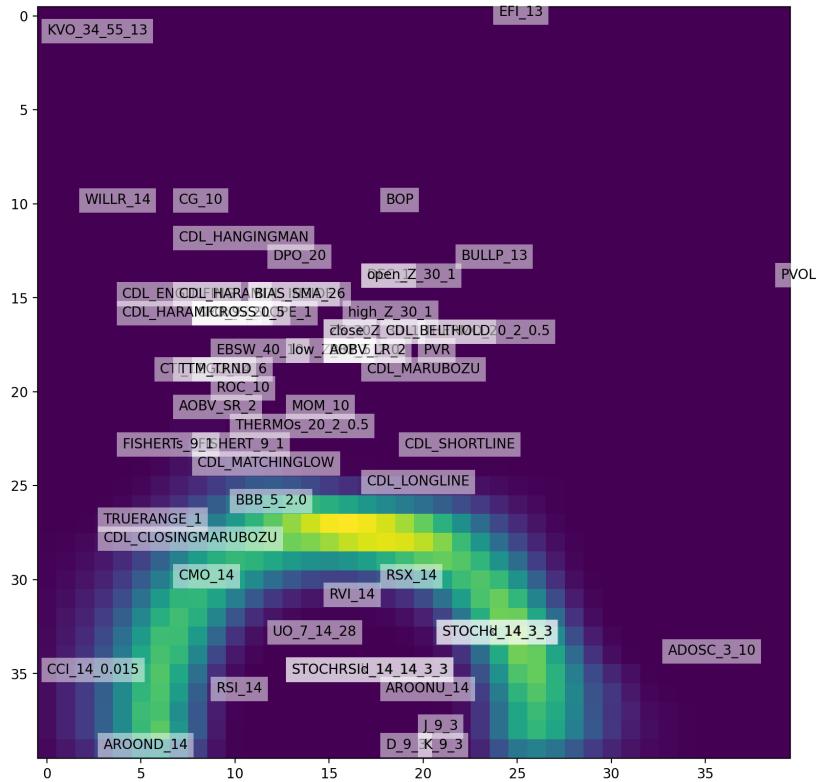
 3.3 SOM

```
In [43]: 1 X = df[['X_mean']]
2 scaler = MinMaxScaler()
3 X = scaler.fit_transform(X)
executed in 6ms, finished 20:09:13 2023-01-23
```

$$Q = 5 \times \sqrt{I}$$

```
In [44]: 1 # Initialize a 15x15 SOM - Thumb rule is 5 * sqrt(n) > 5x5 SOM
2 som = MiniSom(40, 40, 1, learning_rate=0.5, sigma=5, random_seed=2023)
3
4 # Initialize with random weights
5 som.random_weights_init(X)
6
7 # Trains SOM with 10000 iterations
8 som.train_batch(X,10000, verbose=True)
executed in 2.69s, finished 20:09:16 2023-01-23
[ 346 / 10000 ] 3% - 0:00:02 left
[ 665 / 10000 ] 7% - 0:00:02 left
[ 10000 / 10000 ] 100% - 0:00:00 left
quantization error: 7.598641169807632e-12
```

```
In [109]: 1 # Plot SOM
2 plt.figure(figsize=(20, 10))
3 for ix in range(len(X)):
4     winner = som.winner(X[ix])
5     plt.text(winner[0], winner[1], df.Features[ix], bbox=dict(facecolor='white', alpha=0.5, lw=0))
6 plt.imshow(som.distance_map())
7 plt.grid(False)
8 # plt.title('Self Organizing Maps');
executed in 1.32s, finished 21:17:37 2023-01-23
```



```
In [46]: 1 corresponding_listndng_list = []
2 # List the corresponding companies
3 for ix in range(len(X)):
4     winner = som.winner(X[ix])
5     # corresponding_list.append(winner)
6     print(winner[0], winner[1], df['Features'][ix])
executed in 18ms, finished 20:09:18 2023-01-23
```

```
9 18 EBSW_40_10
7 30 CMO_14
6 19 CTI_12
17 14 DEC_1
12 13 DPO_20
2 10 WILLR_14
22 13 BULLP_13
8 23 FISHERT_9_1
4 23 FISHERTS_9_1
15 18 INC_1
8 16 LOGRET_1
13 21 MOM_10
8 16 PCTRET_1
8 19 PGO_14
20 18 PVR
9 20 ROC_10
11 16 SLOPE_1
8 16 SMIO_5_20_5
20 17 THERMO_20_2_0.5
10 22 THERMOS_20_2_0.5
3 27 TRUERANGE_1
7 19 TTM_TRND_6
7 10 CG_10
8 16 CFO_9
15 17 ZS_30
13 18 low_Z_30_1
15 17 close_Z_30_1
18 17 CDL_BELTHOLD
3 28 CDL_CLOSINGMARUBOZU
4 15 CDL_ENGULFING
15 18 BBB_5_2.0
10 26 BBB_5_2.0
3 39 AROOND_14
7 21 AOBV_SR_2
18 10 BOP
7 12 CDL_HANGINGMAN
4 16 CDL_HARAMICROSS
11 15 CDL_INSIDE
15 18 AOBV_LR_2
17 25 CDL_LONGLINE
17 19 CDL_MARUBOZU
8 24 CDL_MATCHINGLOW
19 23 CDL_SHORTLINE
17 14 open_Z_30_1
16 16 high_Z_30_1
7 15 CDL_HARAMI
11 15 BIAS_SMA_26
39 14 PVOL
33 34 ADOSC_3_10
24 0 EFI_13
0 1 KVO_34_55_13
18 30 RSX_14
15 31 RVI_14
20 38 J_9_3
21 33 STOCHK_14_3_3
21 33 STOCHD_14_3_3
13 35 STOCHRSIK_14_14_3_3
13 35 STOCHRSID_14_14_3_3
18 36 AROONU_14
18 39 D_9_3
20 39 K_9_3
12 33 UO_7_14_28
9 36 RSI_14
0 35 CCI_14_0.015

['BEARP_13', 'EBSW_40_10', 'BULLP_13', 'WILLR_14', 'PVR', 'TRUERANGE_1', 'CMO_14', 'AROOND_14', 'BBB_5_2.0', 'PVOL', 'EFI_13', 'Volume', 'KVO_34_55_13', 'AROONU_14', 'CCI_14_0.015']

['EBSW_40_10', 'CMO_14', 'DPO_20', 'WILLR_14', 'BULLP_13', 'PVR', 'TRUERANGE_1', 'ADOSC_3_10', 'EFI_13', 'KVO_34_55_13']

['KVO_34_55_13', 'CCI_14_0.015', 'PVOL', 'EFI_13', 'ADOSC_3_10', 'D_9_3', 'BOP']

In [47]: 1 # selected2 = ['KVO_34_55_13', 'CCI_14_0.015', 'PVOL', 'EFI_13', 'ADOSC_3_10', 'D_9_3', 'BOP']
executed in 2ms, finished 20:09:18 2023-01-23
```

```
In [48]: 1 selected2 = ['KVO_34_55_13', 'EBSW_40_10', 'DPO_20',
2 # 'WILLR_14', 'BULLP_13', 'PVR',
3 # 'TRUERANGE_1', 'EFI_13', 'ADOSC_3_10', 'BOP']
executed in 3ms, finished 20:09:18 2023-01-23
```

```
In [49]: 1 # selected2 = ['EBSW_40_10', 'CMO_14', 'DPO_20',
2 # 'WILLR_14', 'BULLP_13', 'PVR',
3 # 'TRUERANGE_1', 'ADOSC_3_10',
4 # 'EFI_13', 'KVO_34_55_13']
executed in 2ms, finished 20:09:18 2023-01-23
```

<https://github.com/JustGlowing/minisom/blob/master/examples/FeatureSelection.ipynb>

## 4 Exploratory Data Analysis

```
In [50]: 1 X_selected = X_selected[selected2]
executed in 3ms, finished 20:09:18 2023-01-23
```

```
In [51]: 1 X_selected.shape
executed in 5ms, finished 20:09:18 2023-01-23
```

```
Out[51]: (2899, 10)
```

## 4.1 Features Describe

```
In [52]: 1 X_selected.describe().T
executed in 40ms, finished 20:09:18 2023-01-23
Out[52]:
```

	count	mean	std	min	25%	50%	75%	max
KVO_34_55_13	2899.0	-3.152470e+04	7.442190e+06	-5.814442e+07	-4.040862e+06	4.774819e+05	4.347242e+06	2.995325e+07
EBSW_40_10	2899.0	2.145788e-01	8.124232e-01	-9.999737e-01	-7.566179e-01	6.575463e-01	9.683719e-01	9.999898e-01
DPO_20	2899.0	-1.420945e-01	3.917509e+00	-2.789100e+01	-1.595507e+00	3.150635e-02	1.299499e+00	2.668348e+01
WILLR_14	2899.0	-3.583892e+01	3.061310e+01	-1.000000e+02	-5.989283e+01	-2.692308e+01	-8.134081e+00	-0.000000e+00
BULLP_13	2899.0	1.895417e+00	3.997989e+00	-2.681032e+01	3.935218e-01	1.819845e+00	3.505296e+00	2.179554e+01
PVR	2899.0	2.410486e+00	1.016686e+00	1.000000e+00	2.000000e+00	2.000000e+00	3.000000e+00	4.000000e+00
TRUERANGE_1	2899.0	3.254581e+00	3.039621e+00	4.400024e-01	1.389999e+00	2.190002e+00	3.914993e+00	3.196001e+01
EFI_13	2899.0	-2.309790e+07	1.559325e+08	-2.469719e+09	-4.078148e+07	7.807426e+06	3.476777e+07	4.330140e+08
ADOSC_3_10	2899.0	2.980900e+07	7.917299e+07	-4.047211e+08	-1.544632e+07	3.527280e+07	7.614915e+07	3.343292e+08
BOP	2899.0	5.717779e-02	5.265266e-01	-9.838320e-01	-3.874843e-01	7.758349e-02	5.138909e-01	1.000000e+00

## 4.2 Correlation Matrix



```
In [54]: 1 # setting correlation maximum threshold
2 rho_max = 0.80
3
4 # detecting pairwise features with abs(corr) > rho_max
5 idxPairwise = np.where(abs(corrFeatMatrix) >= rho_max)
6 lstPairwise = [
7     (corrFeatMatrix.index[x], corrFeatMatrix.columns[y], round(corrFeatMatrix.iloc[x, y], 3))
8     for x, y in zip(*idxPairwise) if x != y and x < y
9 ]
10 print("::>> Pairwise features with ρ_max= +- 0.80 :")
11 print("::::>>> List/list represents [Feat1, Feat2, Corr]:")
12 lstPairwise
executed in 7ms, finished 20:09:19 2023-01-23
::>> Pairwise features with ρ_max= +- 0.80 :
::::>>> List/list represents [Feat1, Feat2, Corr]:

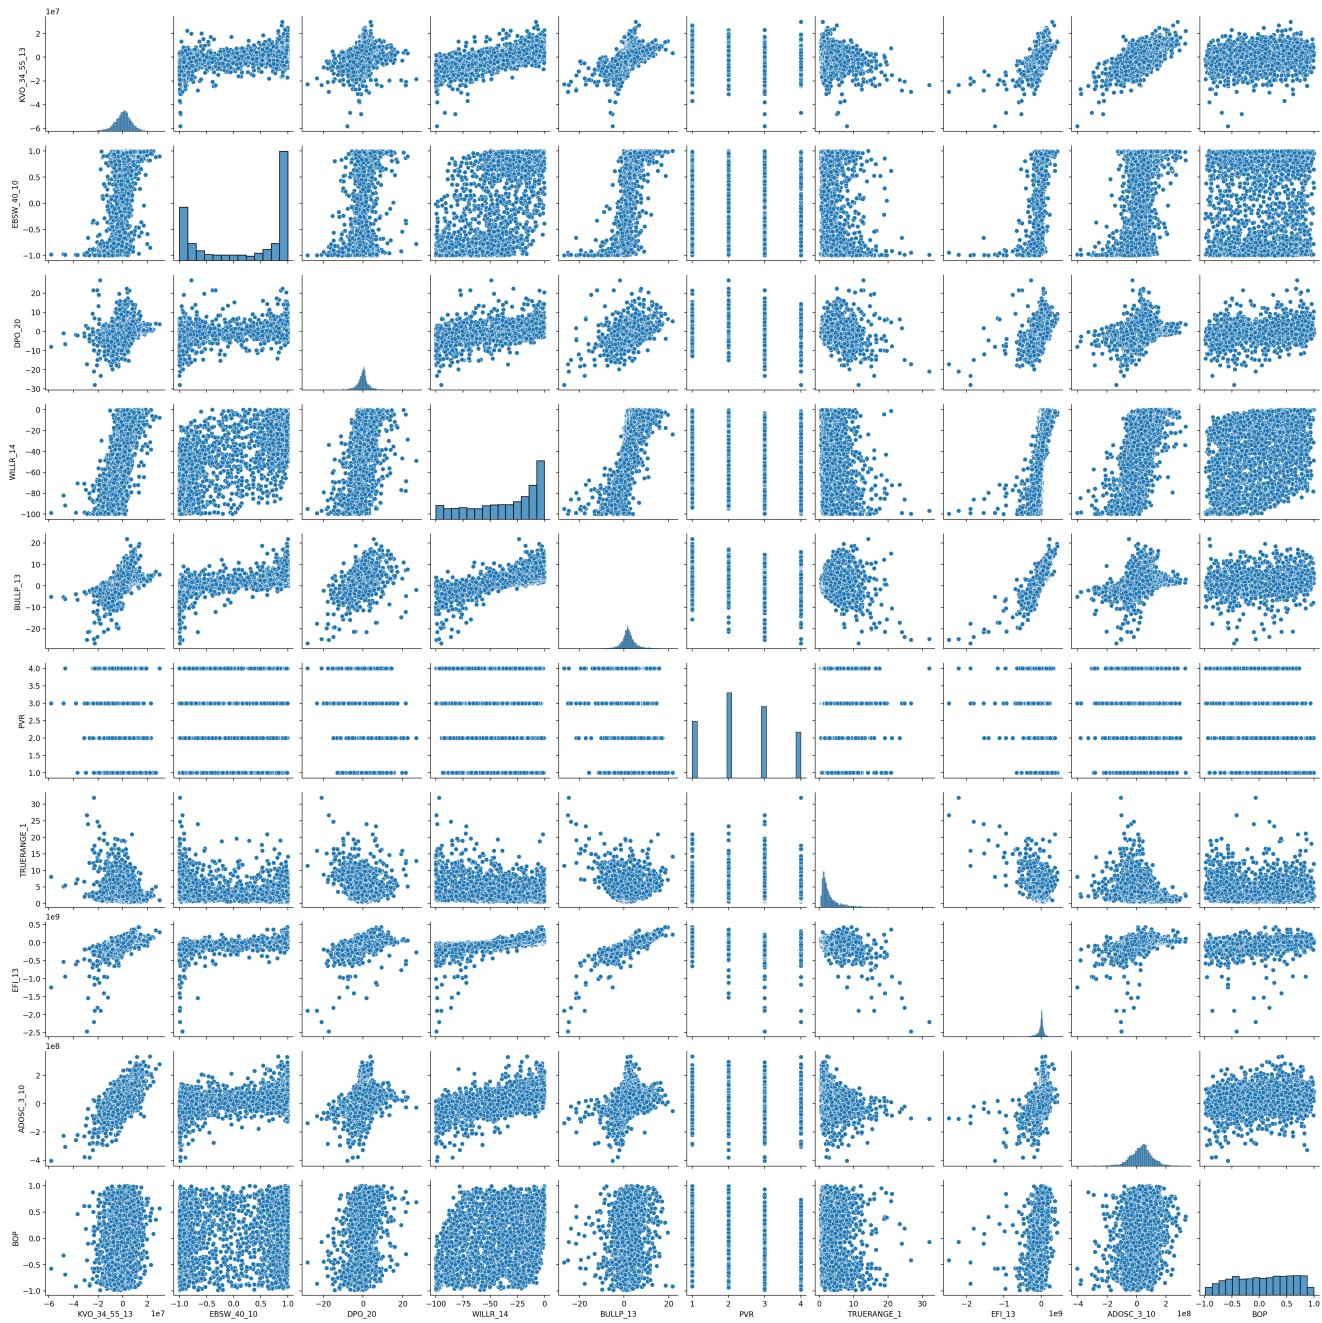
```

Out[54]: []

### 4.3 Pair Plot

In [55]:  
1 sns.pairplot(X\_selected)  
executed in 15.1s, finished 20:09:34 2023-01-23

Out[55]: <seaborn.axisgrid.PairGrid at 0x7fd9742aba00>



## 4.4 XGBoosting

```
In [56]: 1 import shap
2 def XGBoosting_Filter(X_selected, X_train_f, y_train_f, X_test, y_test):
3     # Scale and fit the classifier model
4     xgbcls = XGBClassifier(verbosity = 0, silent=True, random_state=101)
5     xgbcls.fit(X_train_f, y_train_f)
6     # Cross-validation
7     tscv = TimeSeriesSplit(n_splits=5, gap=1)
8     # Hyper parameter optimization
9     param_grid = {'learning_rate': [0.05, 0.10, 0.15, 0.20, 0.25, 0.30],
10                  'max_depth': [3, 4, 5, 6, 8, 10, 12, 15],
11                  'min_child_weight': [1, 3, 5, 7],
12                  'gamma': [0.0, 0.1, 0.2, 0.3, 0.4],
13                  'colsample_bytree': [0.3, 0.4, 0.5, 0.7]}
14
15     # perform random search
16     svm_rdm = RandomizedSearchCV(xgbcls, param_grid, n_iter=100, scoring='f1', cv=tscv, verbose=0)
17     svm_rdm.fit(X_train_f, y_train_f, verbose=0)
18
19     # print best parameter after tuning
20     print("-----")
21     print(svm_rdm.best_params_)
22     # print best score after tuning
23     print("-----")
24     print(svm_rdm.best_score_)
25     print("-----")
26
27     # Refit the XGB Classifier with the best params
28     cls_best = XGBClassifier(**svm_rdm.best_params_)
29
30     cls_best.fit(X_train_f, y_train_f,
31                   eval_set=[(X_train_f, y_train_f), (X_test, y_test)],
32                   eval_metric='logloss',
33                   verbose=True)
34     # Plot feature importance
35     # feature_importance_type = 'gain'
36
37     fig, ax = plt.subplots(figsize=(8, 6))
38     feature_imp = pd.DataFrame({'Importance Score': cls_best.feature_importances_,
39                                 'Features': X_selected.columns}).sort_values(by='Importance Score', ascending=False)
40
41     sns.barplot(x=feature_imp['Importance Score'], y=feature_imp['Features'])
42     ax.set_title('Features Importance')
43     plt.show()
44
45     plot_importance(cls_best, importance_type='gain', show_values=False)
46     plt.show()
47
48     # SHAP
49     explainer = shap.TreeExplainer(cls_best)
50     shap_values = explainer.shap_values(X_test)
51     # future importance summary
52     shap.summary_plot(shap_values, X_test, plot_type="bar")
53     # interpretation plot
54     shap.summary_plot(shap_values, X_test)
```

executed in 2.77s, finished 20:09:37 2023-01-23

## 4.5 Splitting

```
In [57]: 1 # Always keep shuffle = False for financial time series
2 X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.3, shuffle=False)
3
4 # convert to array
5 X_train, X_test, y_train, y_test = np.array(X_train), np.array(X_test), np.array(y_train), np.array(y_test)
6
7 # Output the train and test data size
8 print("Train and Test Size {len(X_train)}, {len(X_test)}")
```

executed in 8ms, finished 20:09:37 2023-01-23

Train and Test Size 2029, 870

```
In [58]: 1 # We scale the data so hasten the LSTM's conversion
2 scaler = MinMaxScaler()
3
4 # Shuffle is set to False because it is a timeseries, and the order matters.
5 X_train, X_test, y_train, y_test = train_test_split(X_selected, y,
6                                                 test_size = 0.30,
7                                                 random_state=0, shuffle = False)
8
9 # Scaling
10 X_scaled_train = scaler.fit_transform(X_train)
11 X_scaled_test = scaler.transform(X_test)
```

executed in 10ms, finished 20:09:37 2023-01-23

```
In [59]: 1 set_seeds(2023)
2 X_scaled_train_df = pd.DataFrame(X_scaled_train)
3 X_scaled_test_df = pd.DataFrame(X_scaled_test)
4 XGBoosting_Filter(X_selected, X_scaled_train_df, y_train,
5                    X_scaled_test_df, y_test)
```

executed in 32.8s, finished 20:10:10 2023-01-23

```
-----  
{'min_child_weight': 7, 'max_depth': 3, 'learning_rate': 0.05, 'gamma': 0.0, 'colsample_bytree': 0.7}  
0.9203230102808355  
-----  
[0] validation_0-logloss:0.65658 validation_1-logloss:0.65313  
[1] validation_0-logloss:0.62371 validation_1-logloss:0.61790  
[2] validation_0-logloss:0.59356 validation_1-logloss:0.58371  
[3] validation_0-logloss:0.57690 validation_1-logloss:0.56552  
[4] validation_0-logloss:0.55099 validation_1-logloss:0.53792  
[5] validation_0-logloss:0.53706 validation_1-logloss:0.52351  
[6] validation_0-logloss:0.51411 validation_1-logloss:0.49722  
[7] validation_0-logloss:0.50220 validation_1-logloss:0.48437  
[8] validation_0-logloss:0.49523 validation_1-logloss:0.47644  
[9] validation_0-logloss:0.47544 validation_1-logloss:0.45377  
[10] validation_0-logloss:0.45714 validation_1-logloss:0.43338  
[11] validation_0-logloss:0.44022 validation_1-logloss:0.41383  
[12] validation_0-logloss:0.43756 validation_1-logloss:0.41082  
[13] validation_0-logloss:0.42204 validation_1-logloss:0.39330  
....
```

In [ ]:

1

## 5 LSTM Model

```
In [60]: 1 # sequence length
2 seqlen = 60
3
4 # number of features
5 numfeat = X_scaled_train.shape[1]
executed in 2ms, finished 20:10:11 2023-01-23
```

```
In [61]: 1 g = TimeseriesGenerator(X_scaled_train, y_train.to_numpy(), length=seqlen)
2 g_ = TimeseriesGenerator(X_scaled_test, y_test.to_numpy(), length=seqlen)
executed in 5ms, finished 20:10:11 2023-01-23
```

## 5.1 Model 1 LSTM Model

```
In [62]: 1 # Create a model
2 def create_model(hu=256, lookback=60, features=numfeat):
3
4     tensorflow.keras.backend.clear_session()
5
6     # instantiate the model
7     model = Sequential()
8     model.add(LSTM(units=hu, input_shape=(lookback, features), activation = 'relu', return_sequences=False, name='LSTM'))
9     model.add(Dense(units=1, name='Output'))           # can also specify linear activation function
10
11    # specify optimizer separately (preferred method)
12    opt = RMSprop(learning_rate=0.001, rho=0.9, epsilon=1e-08)
13    #    opt = Adam(lr=0.001, epsilon=1e-08, decay=0.0)      # adam optimizer seems to perform better for a single lstm
14
15    # model compilation
16    model.compile(optimizer=opt, loss='mse', metrics=['mse'])
17
18    return model
executed in 24ms, finished 20:10:11 2023-01-23
```

```
In [63]: 1 # lstm network
2 model_1 = create_model(hu=32, lookback=seqlen, features=numfeat)
executed in 105ms, finished 20:10:11 2023-01-23
```

```
In [64]: 1 model_1.summary()
executed in 11ms, finished 20:10:11 2023-01-23
```

Model: "sequential"

Layer (type)	Output Shape	Param #
LSTM (LSTM)	(None, 32)	5504
Output (Dense)	(None, 1)	33

Total params: 5,537  
Trainable params: 5,537  
Non-trainable params: 0

```
In [65]: 1 plot_model(model_1, to_file='./img/model_1.png', show_shapes=True, show_layer_names=True)
executed in 443ms, finished 20:10:13 2023-01-23
```

Out[65]:

LSTM_input	input:	[(None, 60, 10)]
InputLayer	output:	[(None, 60, 10)]

↓

LSTM	input:	(None, 60, 10)
LSTM	output:	(None, 32)

↓

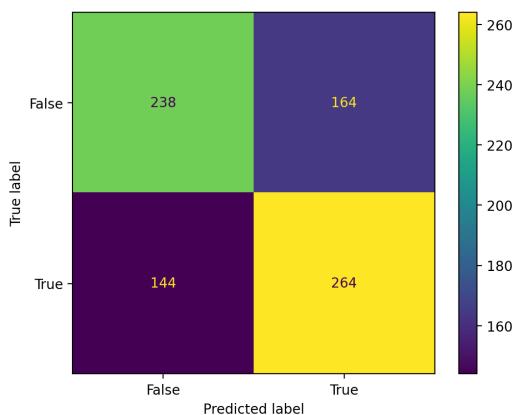
Output	input:	(None, 32)
Dense	output:	(None, 1)

```
In [66]: 1 # Specify callback functions
2 model_path = (results_path / 'model_1.h5').as_posix()
3 logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
4
5 my_callbacks = [
6     EarlyStopping(patience=10, monitor='loss', mode='min', verbose=1, restore_best_weights=True),
7     ModelCheckpoint(filepath=model_path, verbose=1, monitor='loss', save_best_only=True),
8     TensorBoard(log_dir=logdir, histogram_freq=1)
9 ]
executed in 4ms, finished 20:10:13 2023-01-23
```

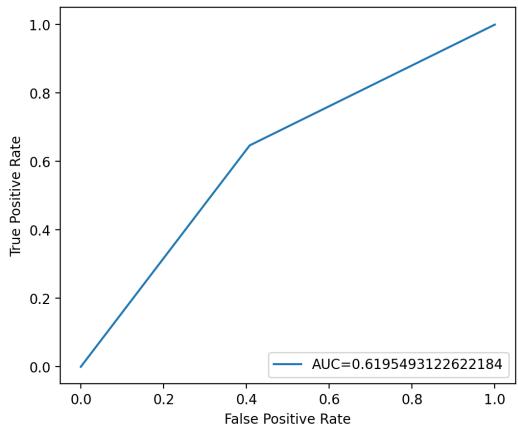
```
In [67]: 1 # Model fitting
2 history = model_1.fit(g,
3                      epochs=500,
4                      verbose=1,
5                      callbacks=my_callbacks,
6                      shuffle=False,
7                      class_weight=class_weight)
executed in 50.7s, finished 20:11:04 2023-01-23
Epoch 8: loss improved from 0.24796 to 0.24744, saving model to results/lstm_time_series/model_1.h5
16/16 [=====] - 0s 17ms/step - loss: 0.2474 - mse: 0.2472
Epoch 9/500
13/16 [=====>.....] - ETA: 0s - loss: 0.2480 - mse: 0.2474
Epoch 9: loss improved from 0.24744 to 0.24692, saving model to results/lstm_time_series/model_1.h5
16/16 [=====] - 0s 18ms/step - loss: 0.2469 - mse: 0.2467
Epoch 10/500
13/16 [=====>.....] - ETA: 0s - loss: 0.2476 - mse: 0.2470
Epoch 10: loss improved from 0.24692 to 0.24642, saving model to results/lstm_time_series/model_1.h5
16/16 [=====] - 0s 17ms/step - loss: 0.2464 - mse: 0.2462
Epoch 11/500
13/16 [=====>.....] - ETA: 0s - loss: 0.2472 - mse: 0.2466
Epoch 11: loss improved from 0.24642 to 0.24591, saving model to results/lstm_time_series/model_1.h5
16/16 [=====] - 0s 17ms/step - loss: 0.2459 - mse: 0.2457
Epoch 12/500
15/16 [=====>.....] - ETA: 0s - loss: 0.2457 - mse: 0.2455
Epoch 12: loss improved from 0.24591 to 0.24541, saving model to results/lstm_time_series/model_1.h5
16/16 [=====] - 0s 21ms/step - loss: 0.2454 - mse: 0.2452
Epoch 13/500
13/16 [=====>.....] - ETA: 0s - loss: 0.2464 - mse: 0.2459
```

```
In [68]: 1 ypred_1 = np.where(model_1.predict(g_, verbose=False) > 0.5, 1, 0)
executed in 224ms, finished 20:11:04 2023-01-23
```

```
In [69]: 1 confusion_matrix = metrics.confusion_matrix(y_test[seqlen:], ypred_1)
2 cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix,
3                                              display_labels = [False, True])
4
5 cm_display.plot()
6 plt.show()
executed in 239ms, finished 20:11:05 2023-01-23
```



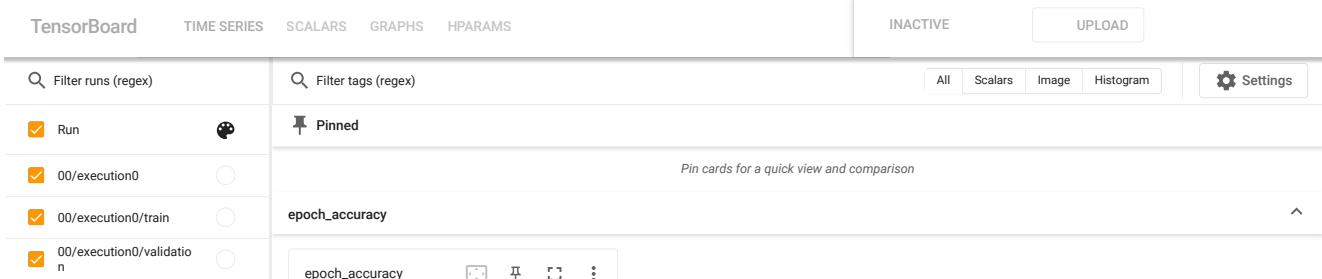
```
In [70]: 1 fpr, tpr, _ = metrics.roc_curve(y_test[seqlen:], ypred_1)
2 auc = metrics.roc_auc_score(y_test[seqlen:], ypred_1)
3
4 #create ROC curve
5 plt.figure(figsize=(6,5))
6 plt.plot(fpr,tpr,label="AUC="+str(auc))
7 plt.ylabel('True Positive Rate')
8 plt.xlabel('False Positive Rate')
9 plt.legend(loc=4)
10 plt.show()
executed in 166ms, finished 20:11:05 2023-01-23
```



```
In [71]: 1 print(classification_report(y_test[seqlen:], ypred_1))
executed in 7ms, finished 20:11:05 2023-01-23
```

	precision	recall	f1-score	support
0	0.62	0.59	0.61	402
1	0.62	0.65	0.63	408
accuracy			0.62	810
macro avg	0.62	0.62	0.62	810
weighted avg	0.62	0.62	0.62	810

```
In [72]: 1 %tensorboard --logdir ./tensorboard/bologs
executed in 10ms, finished 20:11:06 2023-01-23
Reusing TensorBoard on port 6006 (pid 4282), started 1 day, 23:03:04 ago. (Use '!kill 4282' to kill it.)
```



## 5.2 Model 2 Stacked-LSTM Model

gelu + relu + tanh

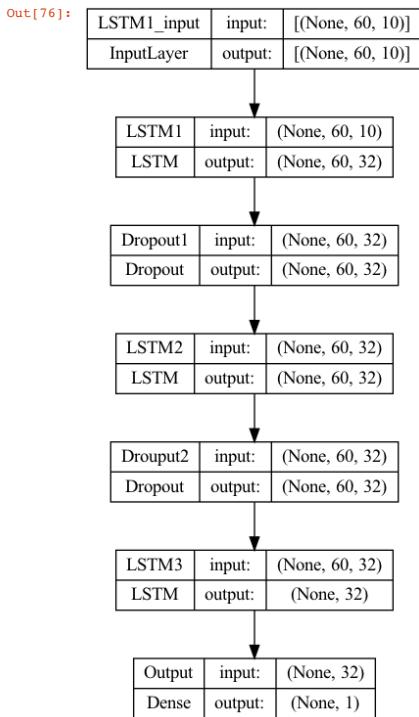
```
In [73]: 1 # Create a model
2 def create_model(hu=256, lookback=60, features=numfeat):
3
4     tensorflow.keras.backend.clear_session()
5
6     # instantiate the model
7     model = Sequential()
8     model.add(LSTM(units=hu, input_shape=(lookback, features), activation = 'relu', return_sequences=True, name='LSTM1'))
9     model.add(Dropout(0.4, name='Dropout1'))
10
11    model.add(LSTM(units=hu, activation = 'relu', return_sequences=True, name='LSTM2'))
12    model.add(Dropout(0.4, name='Dropout2'))
13
14    model.add(LSTM(units=hu, activation = 'selu', return_sequences=False, name='LSTM3'))
15    model.add(Dense(units=1, activation='sigmoid', name='Output')) # can also specify linear activation function
16
17    # specify optimizer separately (preferred method)
18    opt = RMSprop(learning_rate=0.001, rho=0.9, epsilon=1e-08)
19    # opt = Adam(learning_rate=0.001, epsilon=1e-08, decay=0.0) # adam optimizer seems to perform better for a single lstm
20
21    # model compilation
22    model.compile(optimizer=opt, loss='mse', metrics=['mae'])
23
24    return model
executed in 6ms, finished 20:11:06 2023-01-23
```

```
In [74]: 1 # lstm network
2 model_2 = create_model(hu=32, lookback=seqlen, features=numfeat)
executed in 238ms, finished 20:11:07 2023-01-23
```

```
In [75]: 1 model_2.summary()
executed in 14ms, finished 20:11:07 2023-01-23
```

```
Model: "sequential"
+-----+
Layer (type)      Output Shape       Param #
+-----+
LSTM1 (LSTM)      (None, 60, 32)     5504
Dropout1 (Dropout) (None, 60, 32)      0
LSTM2 (LSTM)      (None, 60, 32)     8320
Dropout2 (Dropout) (None, 60, 32)      0
LSTM3 (LSTM)      (None, 32)         8320
Output (Dense)    (None, 1)          33
+-----+
Total params: 22,177
Trainable params: 22,177
Non-trainable params: 0
```

```
In [76]: 1 plot_model(model_2, to_file='./img/model_2.png', show_shapes=True, show_layer_names=True)
executed in 520ms, finished 20:11:08 2023-01-23
```



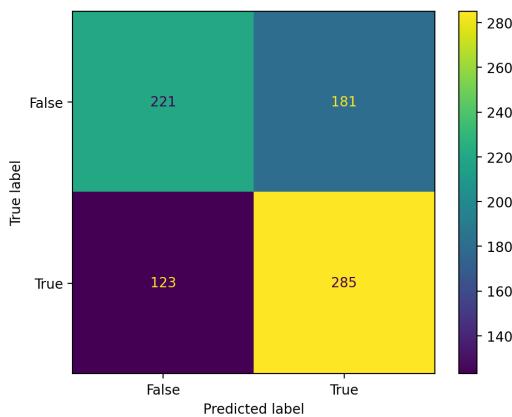
```
In [77]: 1 # Specify callback functions
2 model_path = (results_path / 'model_2.h5').as_posix()
3 logdir = os.path.join("Logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
4
5 my_callbacks = [
6     EarlyStopping(patience=10, monitor='loss', mode='min', verbose=1, restore_best_weights=True),
7     ModelCheckpoint(filepath=model_path, verbose=1, monitor='loss', save_best_only=True),
8     TensorBoard(log_dir=logdir, histogram_freq=1)
9 ]
executed in 6ms, finished 20:11:08 2023-01-23
```

```
In [78]: 1 # Model fitting
2 history = model_2.fit(g,
3     epochs=500,
4     verbose=1,
5     callbacks=my_callbacks,
6     shuffle=False,
7     class_weight=class_weight)
executed in 2m 38s, finished 20:13:46 2023-01-23

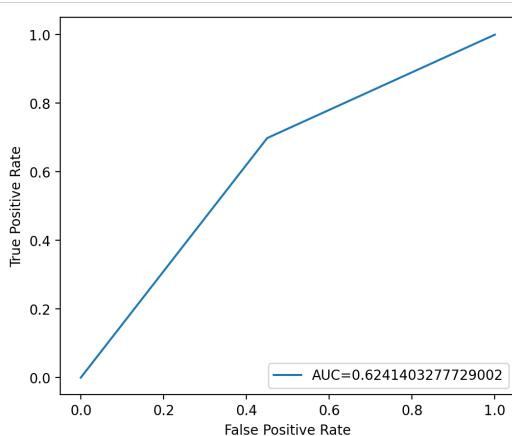
Epoch 1/500
15/16 [=====>..] - ETA: 0s - loss: 0.2515 - mae: 0.5004
Epoch 1: loss improved from inf to 0.25146, saving model to results/lstm_time_series/model_2.h5
16/16 [=====] - 4s 72ms/step - loss: 0.2515 - mae: 0.5004
Epoch 2/500
15/16 [=====>..] - ETA: 0s - loss: 0.2502 - mae: 0.4997
Epoch 2: loss improved from 0.25146 to 0.25012, saving model to results/lstm_time_series/model_2.h5
16/16 [=====] - 1s 75ms/step - loss: 0.2501 - mae: 0.4997
Epoch 3/500
15/16 [=====>..] - ETA: 0s - loss: 0.2501 - mae: 0.4998
Epoch 3: loss improved from 0.25012 to 0.25006, saving model to results/lstm_time_series/model_2.h5
16/16 [=====] - 1s 70ms/step - loss: 0.2501 - mae: 0.4998
Epoch 4/500
15/16 [=====>..] - ETA: 0s - loss: 0.2504 - mae: 0.4999
Epoch 4: loss did not improve from 0.25006
16/16 [=====] - 1s 66ms/step - loss: 0.2503 - mae: 0.4998
Epoch 5/500
15/16 [=====>..] - ETA: 0s - loss: 0.2503 - mae: 0.4998
Epoch 5: loss did not improve from 0.25006
```

```
In [79]: 1 ypred_2 = np.where(model_2.predict(g_, verbose=False) > 0.5, 1, 0
executed in 536ms, finished 20:13:47 2023-01-23
```

```
In [80]: 1 confusion_matrix = metrics.confusion_matrix(y_test[seqlen:], ypred_2)
2 cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix,
3                                              display_labels = [False, True])
4
5 cm_display.plot()
6 plt.show()
```



```
In [81]: 1 fpr, tpr, _ = metrics.roc_curve(y_test[seqlen:], ypred_2)
2 auc = metrics.roc_auc_score(y_test[seqlen:], ypred_2)
3
4 #create ROC curve
5 plt.figure(figsize=(6,5))
6 plt.plot(fpr,tpr,label="AUC="+str(auc))
7 plt.ylabel('True Positive Rate')
8 plt.xlabel('False Positive Rate')
9 plt.legend(loc=4)
10 plt.show()
```



```
In [82]: 1 print(classification_report(y_test[seqlen:], ypred_2))
executed in 9ms, finished 20:13:48 2023-01-23
```

	precision	recall	f1-score	support
0	0.64	0.55	0.59	402
1	0.61	0.70	0.65	408
accuracy			0.62	810
macro avg	0.63	0.62	0.62	810
weighted avg	0.63	0.62	0.62	810

## 6 Hyperparameter Optimization

```

In [83]: 1 def build_model(hp):
2
3     tf.keras.backend.clear_session()
4
5     # instantiate the model
6     model = Sequential()
7
8     # Tune the number of units in the layers
9     hp_units1 = hp.Int('units1', min_value=4, max_value=32, step=4)
10    hp_units2 = hp.Int('units2', min_value=4, max_value=32, step=4)
11    hp_units3 = hp.Int('units3', min_value=4, max_value=32, step=4)
12
13    # Tune the dropout rate
14    hp_dropout1 = hp.Float('Dropout_rate', min_value=0, max_value=0.5, step=0.1)
15    hp_dropout2 = hp.Float('Dropout_rate', min_value=0, max_value=0.5, step=0.1)
16
17    # Tune the learning rate for the optimizer
18    hp_learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])
19
20    # Tune activation functions
21    hp_activation1 = hp.Choice(name = 'activation', values = ['relu', 'gelu', 'elu'], ordered = False)
22    hp_activation2 = hp.Choice(name = 'activation', values = ['relu', 'gelu', 'elu'], ordered = False)
23    hp_activation3 = hp.Choice(name = 'activation', values = ['relu', 'tanh', 'elu'], ordered = False)
24
25    model.add(LSTM(hp_units1, input_shape=(seqlen, numfeat), activation=hp_activation1, return_sequences=True, name='LSTM1'))
26    model.add(Dropout(hp_dropout1, name='Dropout1'))
27
28    model.add(LSTM(hp_units2, activation = hp_activation2, return_sequences=True, name='LSTM2'))
29    model.add(Dropout(hp_dropout2, name='Dropout2'))
30
31    model.add(LSTM(hp_units3, activation = hp_activation3, return_sequences=False, name='LSTM3'))
32
33    model.add(Dense(units=1, activation='sigmoid', name='Output'))
34
35    # specify optimizer separately (preferred method)
36    opt = Adam(learning_rate=hp_learning_rate, epsilon=1e-08)
37
38    # model compilation - 'binary_crossentropy' - 'accuracy' - BinaryAccuracy(name='accuracy', threshold=0.5)
39    model.compile(optimizer=opt,
40                  loss=BinaryCrossentropy(),
41                  metrics=['accuracy',
42                            Precision(),
43                            Recall()])
44
45    return model

```

executed in 7ms, finished 20:13:49 2023-01-23

```

In [84]: 1 # initialize an early stopping callback to prevent the model from
2 # overfitting/spending too much time training with minimal gains
3 callback1 = [EarlyStopping(patience=5, monitor='loss', mode='min', verbose=1, restore_best_weights=True),
4             TensorBoard(log_dir='./tensorboard/rslogs')]
5
6 callback2 = [EarlyStopping(patience=5, monitor='loss', mode='min', verbose=1, restore_best_weights=True),
7             TensorBoard(log_dir='./tensorboard/hblogs')]
8
9 callback3 = [EarlyStopping(patience=5, monitor='loss', mode='min', verbose=1, restore_best_weights=True),
10            TensorBoard(log_dir='./tensorboard/bologs')]

```

executed in 4ms, finished 20:13:49 2023-01-23

## Apple 6.1 Approach 1 : Random Search

[...]

## Apple 6.2 Approach 2 : HyperBand

[...]

## Cat 6.3 Approach 3 : BayesianOptimization

```

In [85]: 1 # BayesianOptimization algorithm from keras tuner
2 botuner = BayesianOptimization(
3     build_model,
4     objectives="val_accuracy",
5     max_trials=15,
6     num_initial_points=2,
7     hyperparameters=None,
8     tune_new_entries=True,
9     allow_new_entries=True,
10    overwrite=True,
11    directory=".\\keras",
12    project_name="botrial")

```

executed in 221ms, finished 20:13:49 2023-01-23

```

In [86]: 1 # launch tuning process
2 botuner.search(g, epochs=50, validation_data=g_, callbacks=callback1, class_weight = class_weight, shuffle=False)

```

executed in 9m 31s, finished 20:23:21 2023-01-23

Trial 15 Complete [00h 00m 56s]  
 val\_accuracy: 0.644444465637207

Best val\_accuracy So Far: 0.6567901372909546  
 Total elapsed time: 00h 09m 31s  
 INFO:tensorflow:Oracle triggered exit

```

In [87]: 1 # display the best hyperparameter values for the model based on the defined objective function
2 best_bohp = botuner.get_best_hyperparameters()[0]
3 print(best_bohp.values)

```

executed in 4ms, finished 20:23:21 2023-01-23

{'units1': 32, 'units2': 32, 'units3': 32, 'Dropout\_rate': 0.2, 'learning\_rate': 0.001, 'activation': 'elu'}

```
In [88]: 1 # display tuning results summary
2 botuner.results_summary()
executed in 5ms, finished 20:23:21 2023-01-23
Dropout_rate: 0.1
learning_rate: 0.001
activation: elu
Score: 0.6370370388031006
Trial summary
Hyperparameters:
units1: 4
units2: 28
units3: 16
Dropout_rate: 0.2
learning_rate: 0.001
activation: elu
Score: 0.5962963104248047
Trial summary
Hyperparameters:
units1: 28
units2: 32
units3: 32
Dropout_rate: 0.0
learning_rate: 0.01
```

```
In [89]: 1 %tensorboard --logdir ./tensorboard/bologs
executed in 10ms, finished 20:23:22 2023-01-23
Reusing TensorBoard on port 6006 (pid 4282), started 1 day, 23:15:20 ago. (Use '!kill 4282' to kill it.)
```

```
In [ ]: 1
```

## 7 Optimal LSTM

```
In [90]: 1 # Create a model
2 def create_model(hu1=32, hu2=32, hu3=32, lookback=60, features=10):
3
4     tensorflow.keras.backend.clear_session()
5
6     # instantiate the model
7     model = Sequential()
8     model.add(LSTM(units=hu1, input_shape=(lookback, features), activation = 'elu', return_sequences=True, name='LSTM1'))
9     model.add(Dropout(0.2, name='Dropout1'))
10
11    model.add(LSTM(units=hu2, activation = 'elu', return_sequences=True, name='LSTM2'))
12    model.add(Dropout(0.2, name='Dropout2'))
13
14    model.add(LSTM(units=hu3, activation = 'elu', return_sequences=False, name='LSTM3'))
15    model.add(Dense(units=1, activation='sigmoid', name='Output'))           # can also specify linear activation function
16
17    # specify optimizer separately (preferred method)
18    opt = RMSprop(learning_rate=0.001, rho=0.9, epsilon=1e-08)
19    # opt = Adam(learning_rate=0.001, epsilon=1e-08, decay=0.0)           # adam optimizer seems to perform better for a single lstm
20
21    # model compilation
22    # model.compile(optimizer=opt, loss='mse', metrics=['mae'])
23    model.compile(optimizer=opt,
24                  loss=BinaryCrossentropy(),
25                  metrics=['accuracy',
26                            Precision(),
27                            Recall()])
28
29    return model
executed in 5ms, finished 20:23:22 2023-01-23
```

```
In [91]: 1 # lstm network
2 model_3 = create_model(hu1=32, hu2=32, hu3=32, lookback=seqlen, features=numfeat)
executed in 304ms, finished 20:23:23 2023-01-23
```

```
In [92]: 1 model_3.summary()
executed in 16ms, finished 20:23:23 2023-01-23
Model: "sequential"

```

Layer (type)	Output Shape	Param #
LSTM1 (LSTM)	(None, 60, 32)	5504
Dropout1 (Dropout)	(None, 60, 32)	0
LSTM2 (LSTM)	(None, 60, 32)	8320
Dropout2 (Dropout)	(None, 60, 32)	0
LSTM3 (LSTM)	(None, 32)	8320
Output (Dense)	(None, 1)	33

---

```
Total params: 22,177
Trainable params: 22,177
Non-trainable params: 0
```

---

```
In [93]: 1 plot_model(model_3, to_file='./img/model_2.png', show_shapes=True, show_layer_names=True)
executed in 549ms, finished 20:23:24 2023-01-23
```

Out[93]:

```

graph TD
    Input[InputLayer] --> LSTM1[LSTM1]
    LSTM1 --> Dropout1[Dropout1]
    Dropout1 --> LSTM2[LSTM2]
    LSTM2 --> Dropout2[Dropout2]
    Dropout2 --> LSTM3[LSTM3]
    LSTM3 --> Output[Output]

```

---

```
In [94]: 1 # Specify callback functions
2 model_path = (results_path / 'model_3.h5').as_posix()
3 logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
4
5 my_callbacks = [
6     EarlyStopping(patience=10, monitor='loss', mode='min', verbose=1, restore_best_weights=True),
7     ModelCheckpoint(filepath=model_path, verbose=1, monitor='loss', save_best_only=True),
8     TensorBoard(log_dir=logdir, histogram_freq=1)
9 ]
executed in 5ms, finished 20:23:24 2023-01-23
```

---

```
In [96]: 1 random.seed(2023)
2 # Model fitting
3 history = model_3.fit(g,
4                      epochs=500,
5                      verbose=1,
6                      callbacks=my_callbacks,
7                      shuffle=False,
8                      class_weight=class_weight)
executed in 35.6s, finished 20:27:25 2023-01-23

```

```

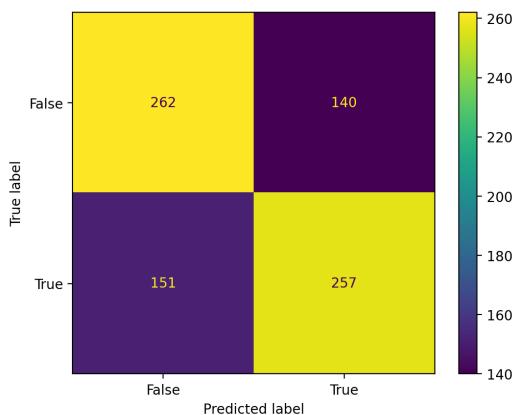
16/16 [=====] - ETA: 0s - loss: 0.6326 - accuracy: 0.6414 - precision: 0.6250 - recall: 0.6143
Epoch 12: loss did not improve from 0.63073
16/16 [=====] - 1s 83ms/step - loss: 0.6326 - accuracy: 0.6414 - precision: 0.6250 - recall: 0.6143
Epoch 13/500
16/16 [=====] - ETA: 0s - loss: 0.6298 - accuracy: 0.6425 - precision: 0.6245 - recall: 0.6218
Epoch 13: loss improved from 0.63073 to 0.62981, saving model to results/lstm_time_series/model_3.h5
16/16 [=====] - 1s 74ms/step - loss: 0.6298 - accuracy: 0.6425 - precision: 0.6245 - recall: 0.6218
Epoch 14/500
15/16 [=====] - ETA: 0s - loss: 0.6333 - accuracy: 0.6375 - precision: 0.6195 - recall: 0.6182
Epoch 14: loss did not improve from 0.62981
16/16 [=====] - 1s 63ms/step - loss: 0.6329 - accuracy: 0.6374 - precision: 0.6194 - recall: 0.6154
Epoch 15/500
15/16 [=====] - ETA: 0s - loss: 0.6432 - accuracy: 0.6307 - precision: 0.6089 - recall: 0.6269
Epoch 15: loss did not improve from 0.62981
16/16 [=====] - 1s 62ms/step - loss: 0.6418 - accuracy: 0.6318 - precision: 0.6100 - recall: 0.6250
Epoch 16/500
16/16 [=====] - ETA: 0s - loss: 0.6290 - accuracy: 0.6486 - precision: 0.6298 - recall: 0.6325
Epoch 16: loss improved from 0.62981 to 0.62905, saving model to results/lstm_time_series/model_3.h5
16/16 [=====] - 1s 83ms/step - loss: 0.6290 - accuracy: 0.6486 - precision: 0.6298 - recall: 0.6325
Epoch 17/500

```

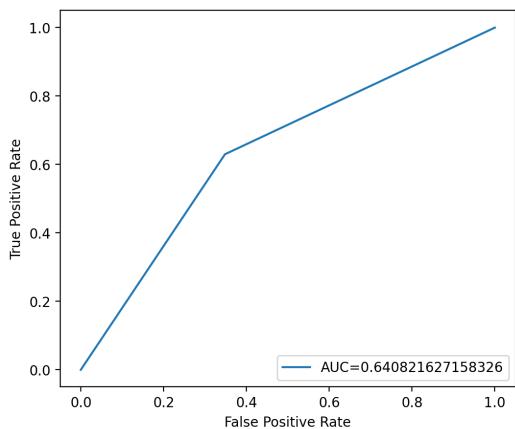
---

```
In [97]: 1 ypred_3 = np.where(model_3.predict(g_, verbose=False) > 0.5, 1, 0)
executed in 621ms, finished 20:27:26 2023-01-23
```

```
In [98]:  
1 confusion_matrix = metrics.confusion_matrix(y_test[seqlen:], ypred_3)  
2 cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix,  
3                                              display_labels = [False, True])  
4  
5 cm_display.plot()  
6 plt.show()  
executed in 252ms, finished 20:27:27 2023-01-23
```



```
In [99]:  
1 fpr, tpr, _ = metrics.roc_curve(y_test[seqlen:], ypred_3)  
2 auc = metrics.roc_auc_score(y_test[seqlen:], ypred_3)  
3  
4 #create ROC curve  
5 plt.figure(figsize=(6,5))  
6 plt.plot(fpr,tpr,label="AUC="+str(auc))  
7 plt.ylabel('True Positive Rate')  
8 plt.xlabel('False Positive Rate')  
9 plt.legend(loc=4)  
10 plt.show()  
executed in 167ms, finished 20:27:28 2023-01-23
```



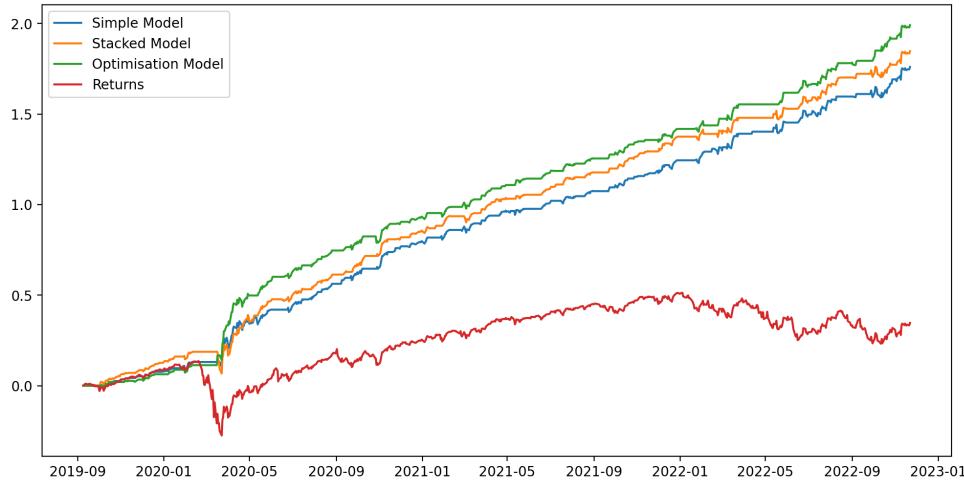
```
In [100]:  
1 print(classification_report(y_test[seqlen:], ypred_3))  
executed in 7ms, finished 20:27:28 2023-01-23
```

	precision	recall	f1-score	support
0	0.63	0.65	0.64	402
1	0.65	0.63	0.64	408
accuracy			0.64	810
macro avg	0.64	0.64	0.64	810
weighted avg	0.64	0.64	0.64	810

## 8 Trading Strategy

## 8.1 Long-Only Strategy

```
In [101]:  
1 bk_lo = stock_data[['Returns']].loc[X_test.iloc[60:,:].index]  
2 bk_lo['Sign_1'] = ypred_1  
3 bk_lo['Sign_2'] = ypred_2  
4 bk_lo['Sign_3'] = ypred_3  
5 bk_lo['Model_1'] = bk_lo['Returns'] * bk_lo['Sign_1'].fillna(0)  
6 bk_lo['Model_2'] = bk_lo['Returns'] * bk_lo['Sign_2'].fillna(0)  
7 bk_lo['Model_3'] = bk_lo['Returns'] * bk_lo['Sign_3'].fillna(0)  
8  
9 plt.figure(figsize=(12,6))  
10 plt.plot(np.cumsum(bk_lo['Model_1']), label='Simple Model')  
11 plt.plot(np.cumsum(bk_lo['Model_2']), label='Stacked Model')  
12 plt.plot(np.cumsum(bk_lo['Model_3']), label='Optimisation Model')  
13 plt.plot(np.cumsum(bk_lo['Returns']), label='Returns')  
14 plt.legend()  
15 plt.show()  
executed in 323ms, finished 20:27:29 2023-01-23
```

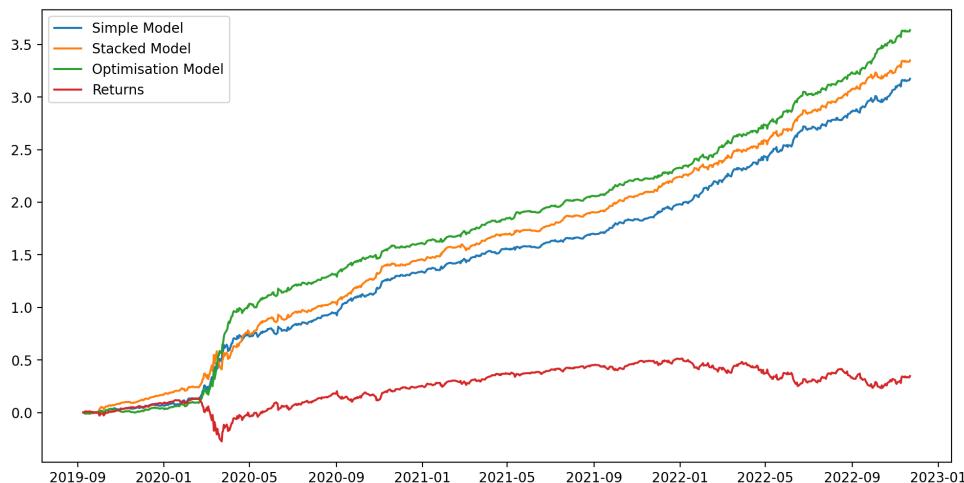


```
In [102]:  
1 print('Simple LSTM Return:', round(np.cumsum(bk_lo['Model_1'])[-1]*100, 2), '%')  
2 print('Stacked LSTM Return:', round(np.cumsum(bk_lo['Model_2'])[-1]*100, 2), '%')  
3 print('Optimal LSTM Return:', round(np.cumsum(bk_lo['Model_3'])[-1]*100, 2), '%')  
4 print('Holding Return:', round(np.cumsum(bk_lo['Returns'])[-1]*100, 2), '%')  
executed in 6ms, finished 20:27:30 2023-01-23
```

Simple LSTM Return: 176.01 %  
Stacked LSTM Return: 184.74 %  
Optimal LSTM Return: 199.13 %  
Holding Return: 34.67 %

## 8.2 Long-Short Strategy

```
In [103]:  
1 bk_ls = stock_data[['Returns']].loc[X_test.iloc[60:,:].index]  
2 bk_ls['Sign_1'] = bk_lo['Sign_1'].apply(lambda sig: 1 if sig == 1 else (-1 if sig == 0 else sig))  
3 bk_ls['Sign_2'] = bk_lo['Sign_2'].apply(lambda sig: 1 if sig == 1 else (-1 if sig == 0 else sig))  
4 bk_ls['Sign_3'] = bk_lo['Sign_3'].apply(lambda sig: 1 if sig == 1 else (-1 if sig == 0 else sig))  
5  
6 bk_ls['Model_1'] = bk_ls['Returns'] * bk_ls['Sign_1'].fillna(0)  
7 bk_ls['Model_2'] = bk_ls['Returns'] * bk_ls['Sign_2'].fillna(0)  
8 bk_ls['Model_3'] = bk_ls['Returns'] * bk_ls['Sign_3'].fillna(0)  
9  
10 plt.figure(figsize=(12,6))  
11 plt.plot(np.cumsum(bk_ls['Model_1']), label='Simple Model')  
12 plt.plot(np.cumsum(bk_ls['Model_2']), label='Stacked Model')  
13 plt.plot(np.cumsum(bk_ls['Model_3']), label='Optimisation Model')  
14 plt.plot(np.cumsum(bk_ls['Returns']), label='Returns')  
15 plt.legend()  
16 plt.show()  
executed in 301ms, finished 20:27:30 2023-01-23
```



```
In [104]:  
1 print('Simple LSTM Return:', round(np.cumsum(bk_ls['Model_1'])[-1]*100, 2), '%')  
2 print('Stacked LSTM Return:', round(np.cumsum(bk_ls['Model_2'])[-1]*100, 2), '%')  
3 print('Optimal LSTM Return:', round(np.cumsum(bk_ls['Model_3'])[-1]*100, 2), '%')  
4 print('Holding Return:', round(np.cumsum(bk_ls['Returns'])[-1]*100, 2), '%')  
  
executed in 8ms, finished 20:27:31 2023-01-23  
Simple LSTM Return: 317.34 %  
Stacked LSTM Return: 334.8 %  
Optimal LSTM Return: 363.58 %  
Holding Return: 34.67 %
```

In [ ]: 1

In [ ]: 1

## References

- Hilpisch, Y. (2020), *Artificial Intelligence in Finance*, O'Reilly Media.
- Kannan Singaravelu, C. (2022a), *Feature Engineering : Extraction & Selection*, Machine Learning Workshop Notes.
- Kannan Singaravelu, C. (2022b), *KMeans Clustering*, CQF Python Lab Notes.
- Kannan Singaravelu, C. (2022c), *Neural Networks: Deep Sequence Modeling*, Machine Learning Workshop Notes.
- Kannan Singaravelu, C. (2022d), *Self Organizing Maps*, CQF Python Lab Notes.
- Kursa, M. B. & Rudnicki, W. R. (2010), 'Feature selection with the boruta package', *Journal of statistical software* **36**, 1–13.
- Lau, K. W., Yin, H. & Hubbard, S. (2006), 'Kernel self-organising maps for classification', *Neurocomputing* **69**(16-18), 2033–2040.
- Xu, D. & Tian, Y. (2015), 'A comprehensive survey of clustering algorithms', *Annals of Data Science* **2**(2), 165–193.
- Zhou, H. B. & Gao, J. T. (2014), Automatic method for determining cluster number based on silhouette coefficient, in 'Advanced materials research', Vol. 951, Trans Tech Publ, pp. 227–230.
- Zou, Z. & Qu, Z. (2020), 'Using lstm in stock prediction and quantitative trading', *CS230: Deep Learning, Winter* pp. 1–6.